

MSiA 400 Lab Advanced Regression with R

Oct 27, 2014

Young-Woong Park

Variable Selection in Multiple Regression

- Goal: choose the candidate variables to obtain a regression model that contains the “best” subset of regressor variables
- Methods
 - All possible regression
 - Stepwise regression
- Criteria
 - $C_p = \frac{SSE_p}{MSE_p} + 2p - n$
 - $AIC = n \log \frac{SSE_p}{n} + 2p$
 - Adjusted $r^2 = 1 - \frac{SSR/(n-p-1)}{SST/(n-1)}$

Data Set for Multiple Regression

- Data set: Wine quality (white wine)
 - Number of observations: 4898
 - Number of attributes: 11 + output attribute
 - Input attributes: fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, alcohol
 - Output attribute: quality (score between 0 and 10)

Ref: P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis., Modeling wine preferences by data mining from physicochemical properties, Decision Support Systems, Elsevier, 47(4):547-553. ISSN: 0167-9236.

- Read the data set

```
> wine <- read.delim("../whitewine.txt");  
> y = wine[,1];  
> x = wine[,2:length(wine[1,])];
```

Stepwise Regression with AIC

- Stepwise regression

```
> library(MASS)
> reg = lm(y~., data=x)
> reg.step = stepAIC(object=reg, direction="both")
```

```
Step:  AIC=-2792.2
b ~ x1 + x2 + x4 + x6 + x7 + x8 + x9 + x10 + x11
```

“forward”

“backward”

	Df	Sum of Sq	RSS	AIC
- x7	1	0.320	2758.8	-2793.6
<none>			2758.5	-2792.2
+ x5	1	0.105	2758.4	-2790.4
+ x3	1	0.019	2758.4	-2790.2
- x1	1	6.157	2764.6	-2783.3
- x6	1	11.036	2769.5	-2774.7
- x10	1	22.570	2781.0	-2754.3
- x9	1	25.297	2783.8	-2749.5
- x11	1	36.536	2795.0	-2729.8
- x8	1	36.823	2795.3	-2729.2
- x4	1	70.134	2828.6	-2671.2
- x2	1	158.543	2917.0	-2520.5

```
Step:  AIC=-2793.63
b ~ x1 + x2 + x4 + x6 + x8 + x9 + x10 + x11
```

Stepwise Regression with AIC (Cont.)

- Displaying the summary

```
> summary(reg.step)
```

```
Call:
```

```
lm(formula = b ~ x1 + x2 + x4 + x6 + x8 + x9 + x10 + x11, data = a)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-3.8246	-0.4938	-0.0396	0.4660	3.1208

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.541e+02	1.810e+01	8.514	< 2e-16	***
x1	6.810e-02	2.043e-02	3.333	0.000864	***
x2	-1.888e+00	1.095e-01	-17.242	< 2e-16	***
x4	8.285e-02	7.287e-03	11.370	< 2e-16	***
x6	3.349e-03	6.766e-04	4.950	7.67e-07	***
x8	-1.543e+02	1.834e+01	-8.411	< 2e-16	***
x9	6.942e-01	1.034e-01	6.717	2.07e-11	***
x10	6.285e-01	9.997e-02	6.287	3.52e-10	***
x11	1.932e-01	2.408e-02	8.021	1.31e-15	***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.7512 on 4889 degrees of freedom
```

```
Multiple R-squared:  0.2818, Adjusted R-squared:  0.2806
```

```
F-statistic: 239.7 on 8 and 4889 DF,  p-value: < 2.2e-16
```

Stepwise Regression with AIC (Cont.)

- Calculating / Referencing Statistics

> formula(reg.step);	# print the formula of the model
> AIC(reg. step);	# print AIC value of the model
> summary(reg. step)\$r.squared;	# print r^2 value of the model
> summary(reg. step)\$adj.r.squared;	# print adjusted r^2 value
> e = resid(reg. step);	# define residuals
> SSE = sum(e^2);	# calculate Sum of Squared errors
> SAE = sum(abs(e));	# calculate Sum of Absolute errors

Variable Selection Using Package

- Load package leaps

```
> library(leaps)
```

- Finding the best subset

- For number of variables $p=1,2,\dots,nvmax$,
Find nbest best subsets with cardinality p

```
> reg.exh = regsubsets(x,y, nbest=1, nvmax=length(y), method="exhaustive");  
> summary(reg.exh)
```

```
1 subsets of each size up to 11  
Selection Algorithm: exhaustive  
      x1  x2  x3  x4  x5  x6  x7  x8  x9  x10 x11  
1  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
2  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
3  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
4  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
5  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
6  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
7  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
8  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
9  ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
10 ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "  
11 ( 1 )  " " " " " " " " " " " " " " " " " " " " " " " " " " " "
```

Variable Selection Using Package (Cont.)

- Calculating / Referencing Statistics

```
> summary(reg.exh)$which  
> summary(reg.exh)$cp  
> summary(reg.exh)$adjr2  
> cbind(summary(reg.exh)$which, summary(reg.exh)$cp, summary(reg.exh)$adjr2)
```

	(Intercept)	x1	x2	x3	x4	x5	x6	x7	x8	x9	x10	x11			
1	1	0	0	0	0	0	0	0	0	0	0	0	1	618.935028	0.1895598
2	1	0	1	0	0	0	0	0	0	0	0	0	1	277.304045	0.2399208
3	1	0	1	0	1	0	0	0	0	0	0	0	1	154.828971	0.2580716
4	1	0	1	0	1	0	1	0	0	0	0	0	1	119.625443	0.2633925
5	1	0	1	0	1	0	0	0	1	1	0	0	1	73.461400	0.2703282
6	1	0	1	0	1	0	0	0	1	1	1	1	1	37.464938	0.2757705
7	1	0	1	0	1	0	1	0	1	1	1	1	1	15.911941	0.2790891
8	1	1	1	0	1	0	1	0	1	1	1	1	1	6.805571	0.2805767
9	1	1	1	0	1	0	1	1	1	1	1	1	1	8.238314	0.2805130
10	1	1	1	0	1	1	1	1	1	1	1	1	1	10.053204	0.2803931
11	1	1	1	1	1	1	1	1	1	1	1	1	1	12.000000	0.2802536

Variable Selection Using Package (Cont.)


- Optimizing various criteria

```
> leaps(x,y,nbest=1,method="Cp")  
> leaps(x,y,nbest=1,method="adjr2")
```

```
> leaps(x,y, nbest=1, method="Cp");  
$which  
      1      2      3      4      5      6      7      8      9      A      B  
1 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE  
2 FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE  
3 FALSE TRUE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE TRUE  
4 FALSE TRUE FALSE TRUE FALSE TRUE FALSE FALSE FALSE FALSE TRUE  
5 FALSE TRUE FALSE TRUE FALSE FALSE FALSE TRUE TRUE FALSE TRUE  
6 FALSE TRUE FALSE TRUE FALSE FALSE FALSE TRUE TRUE TRUE TRUE  
7 FALSE TRUE FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE  
8 TRUE TRUE FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE  
9 TRUE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE  
10 TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE  
11 TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```


```
> leaps(x,y, nbest=1, method=c("adjr2"));  
$which  
      1      2      3      4      5      6      7      8      9      A      B  
1 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE  
2 FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE  
3 FALSE TRUE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE TRUE  
4 FALSE TRUE FALSE TRUE FALSE TRUE FALSE FALSE FALSE FALSE TRUE  
5 FALSE TRUE FALSE TRUE FALSE FALSE FALSE TRUE TRUE FALSE TRUE  
6 FALSE TRUE FALSE TRUE FALSE FALSE FALSE TRUE TRUE TRUE TRUE  
7 FALSE TRUE FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE  
8 TRUE TRUE FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE  
9 TRUE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE  
10 TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE  
11 TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

Piecewise Regression



Is simple linear regression working?

Piecewise Regression




Is simple linear regression working?

Piecewise Regression

- Load height2.txt

```
> height2 <- read.delim("../height2.txt");  
> mht = height2[,1];  
> fht = height2[,2];
```



Piecewise Regression (Cont.)

- Let us set male height 1.73 as a breakpoint

```
> reg.seg = lm(fht ~ (mht<1.73)*mht )
```

Call:

```
lm(formula = fht ~ (mht < 1.73) * mht)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.057425	-0.010236	-0.000560	0.009237	0.055954

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.2517	0.1867	1.348	0.1822
mht < 1.73TRUE	-0.5669	0.2365	-2.397	0.0193 *
mht	0.7857	0.1051	7.478	2.1e-10 ***
mht < 1.73TRUE:mht	0.3359	0.1362	2.467	0.0162 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.01792 on 67 degrees of freedom

Multiple R-squared: 0.9027, Adjusted R-squared: 0.8984

F-statistic: 207.2 on 3 and 67 DF, p-value: < 2.2e-16

Piecewise Regression (Cont.)

- Interpreting the result

```
Call:
lm(formula = fht ~ (mht < 1.73) * mht)

Residuals:
    Min       1Q   Median       3Q      Max
-0.057425 -0.010236 -0.000560  0.009237  0.055954

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    0.2517     0.1867    1.348   0.1822
mht < 1.73TRUE -0.5669     0.2365   -2.397   0.0193 *
mht             0.7857     0.1051    7.478  2.1e-10 ***
mht < 1.73TRUE:mht 0.3359     0.1362    2.467   0.0162 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.01792 on 67 degrees of freedom
Multiple R-squared:  0.9027, Adjusted R-squared:  0.8984
F-statistic: 207.2 on 3 and 67 DF, p-value: < 2.2e-16
```

$$fht = 0.2517 - 0.5669 * 1_{mht < 1.73} + 0.7857mht + 0.3359 * 1_{mht < 1.73} * mht$$

Note

- For an observation with male height ≥ 1.73 , we have

$$\begin{aligned} fht &= 0.2517 - 0.5669 * 0 + 0.7857mht + 0.3359 * 0 * mht \\ &= 0.2517 + 0.7857mht \end{aligned}$$

- For an observation with male height < 1.73 , we have

$$\begin{aligned} fht &= 0.2517 - 0.5669 * 1 + 0.7857mht + 0.3359 * 1 * mht \\ &= -0.3152 + 1.1216mht \end{aligned}$$

Piecewise Regression Using Package

- Load package segmented

```
> library(segmented)
```

- Let us **guess** male height 1.73 as a breakpoint

```
> reg.ht = lm(fht ~ mht)
> reg.seg1 = segmented(reg.ht, seg.Z = ~mht, psi=1.73)
```

```
***Regression Model with Segmented Relationship(s)***
```

```
Call:
```

```
segmented.lm(obj = reg.ht, seg.Z = ~mht, psi = 1.6)
```

```
Estimated Break-Point(s):
```

```
Est. St.Err
1.66400 0.01238
```

$$fht = -0.8743 + 1.4642 * mht - 0.6994 * 1_{mht < 1.664} * mht$$

```
t value for the gap-variable(s) v: 0
```

```
Meaningful coefficients of the linear terms:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.8743	0.3392	-2.577	0.0122 *
mht	1.4642	0.2069	7.077	1.1e-09 ***
U1.mht	-0.6994	0.2141	-3.267	NA


```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.01761 on 67 degrees of freedom
```


```
Multiple R-Squared: 0.9061, Adjusted R-squared: 0.9019
```

Piecewise Regression Using Package(Cont.)

- Plotting the result



```
> plot(ht)  
> plot(reg.seg1, add=T)
```



```
> plot(ht)  
> abline(coef(reg.ht))
```


Piecewise Regression Using Package(Cont.)

- Piecewise regression with multiple breakpoints?

```
> reg.ht = lm(fht ~ mht)
> reg.seg1 = segmented(reg.ht, seg.Z = ~mht, psi=1.73)
> reg.seg2 = segmented(reg.ht, seg.Z = ~mht, psi=c(1.65,1.73))
```