

MSIA 401 - Hw6

Steven Lin

Setup

```
# My PC
main = "C:/Users/Steven/Documents/Academics/3_Graduate School/2014-2015
~ NU/"

# Aginity main = '\\\\nas1/labuser169'

course = "MSIA_401_Statistical Methods for Data Mining"
datafolder = "Data"
setwd(file.path(main, course, datafolder))

opts_knit$set(root.dir = getwd())
```

Problem 1

```
# Import data
filename = "P219.txt"
mydata = read.table(filename, header = T)
```

Part a

```
library(car)
fit = lm(H ~ P, data = mydata)
summary(fit)
```

```
##
## Call:
## lm(formula = H ~ P, data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.008368 -0.002133  0.000525  0.002557  0.008075
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.06088    0.01042   -5.85  5.9e-06 ***
## P             0.07141    0.00423   16.87  1.9e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.00408 on 23 degrees of freedom
## Multiple R-squared:  0.925,    Adjusted R-squared:  0.922
## F-statistic: 285 on 1 and 23 DF,  p-value: 1.91e-14
```

```
dw_positive = durbinwatsonTest(fit, alternative = "positive", data =
mydata)
dw_2sided = durbinwatsonTest(fit, alternative = "two.sided", data =
mydata)
dw_positive
```

```
## lag Autocorrelation D-W Statistic p-value
## 1 0.6511 0.6208 0
## Alternative hypothesis: rho > 0
```

```
dw_2sided
```

```
## lag Autocorrelation D-W Statistic p-value
## 1 0.6511 0.6208 0
## Alternative hypothesis: rho != 0
```

```
# alternatives library(car)
# durbinwatsonTest(fit_lagged,alternative='two.sided')
# durbinwatsonTest(fit_lagged,alternative='positive')

library(lmtest)
```

```
## Loading required package: zoo
##
## Attaching package: 'zoo'
##
## The following object(s) are masked from 'package:base':
##
## as.Date, as.Date.numeric
```

```
dwtest(fit, alternative = "greater")
```

```
##
## Durbin-Watson test
##
## data: fit
## DW = 0.6208, p-value = 6.645e-06
## alternative hypothesis: true autocorrelation is greater than 0
```

```
dwtest(fit, alternative = "two.sided")
```

```
##
## Durbin-Watson test
##
## data: fit
## DW = 0.6208, p-value = 1.329e-05
## alternative hypothesis: true autocorrelation is not 0
```

```
# Durbin-Watson statistic
dw_positive$dw
```

```
## [1] 0.6208
```

Evidence of autocorrelation is indicated by the deviation of d from 2

Durbin watson: d= 0.6208

H0: correlation = 0, H1: correlation > 0 From talbe A.6, with n = 25, p = 1, and significance level 0.05, dL = 1.29, dU = 1.45. Since d = 0.62 < dL = 1.29, reject null, conclude that value of d is significant at 0.05 level, showing that positive autocorrelation is present. Similar conclusion is reached for two-sided hypothesis (H1: correlation diff 0)


Part b

```
# compute standard residuals
fit_stdres = rstandard(fit)
index_res = seq(1:length(fit_stdres))


# add standard residuals and color code +/- residuals
mydata$stdres = rstandard(fit)
mydata$res_sign[mydata$stdres > 0] = 1
mydata$res_sign[mydata$stdres < 0] = -1
mydata$color[mydata$stdres > 0] = "red"
mydata$color[mydata$stdres < 0] = "blue"

# plot results

plot(index_res, fit_stdres, ylab = "Standardized Residuals", xlab =
"Index",
      col = mydata$color, pch = 16)
abline(0, 0, lty = 2)
lines(index_res, fit_stdres)
```



```
plot(index_res, mydata$res_sign, ylab = "Sign Residuals", xlab =  
"Index", col = mydata$color,  
      pch = 16)  
abline(0, 0, lty = 2)
```



```
# n1 = # of + res, n2 = # of - res
n1 = length(which(fit_stdres > 0))
n2 = length(which(fit_stdres < 0))

# expected value and standard deviation of number of runs
mu = 2 * n1 * n2 / (n1 + n2) + 1
sigma = sqrt(2 * n1 * n2 * (2 * n1 * n2 - n1 - n2) / ((n1 + n2 - 1) * (n1 + n2)^2))

# number of runs
mydata$res_sign_lag = c(mydata$res_sign[1],
mydata$res_sign[1:length(mydata$res_sign) -
1])

mydata$res_sign_change = mydata$res_sign != mydata$res_sign_lag
n_sign_changes = sum(mydata$res_sign_change)
n_runs = n_sign_changes + 1

# mu sigma n_runs
#
# mu-n_runs (mu-n_runs)/sigma
```

- Observed number of runs: 6.
- Expected number of runs: 13.32.
- Standard deviation: 2.4106.

The deviation of 7.32 from the expected number of runs is more than triple the standard deviation, indicating a significant departure from randomness.

```
# Using a statistical test:
z = (n_runs - mu)/sigma

# Compute critical value (two-sided)
z_crit = qnorm(0.05/2, lower.tail = FALSE)
z
```

```
## [1] -3.037
```

```
z_crit
```

```
## [1] 1.96
```

```
abs(z) > z_crit
```

```
## [1] TRUE
```

Reject null hypothesis that sequence is random and conclude that there is autocorrelation present

```
# Compute critical value (one-sided)
z_crit = qnorm(0.05, lower.tail = TRUE)
z
```

```
## [1] -3.037
```

```
z_crit
```

```
## [1] -1.645
```

```
z < z_crit
```

```
## [1] TRUE
```

Positive autocorrelation is manifested by Small values of number of runs and hence small negative values of Z. Reject null hypothesis that sequence is random and conclude that there is positive autocorrelation present

```
# use package source:
# http://www.itl.nist.gov/div898/handbook/eda/section3/eda35d.htm
library(lawstat)


# two.sided
run_test = runs.test(fit_stdres, plot.it = TRUE, alternative =
"two.sided")
run_test
```

```
##
## Runs Test - Two sided
##
## data: fit_stdres
## Standardized Runs Statistic = -3.062, p-value = 0.002203
```

```
# Compute critical value.
qnorm(0.05/2, lower.tail = FALSE)
```

```
## [1] 1.96
```

```
# positive.correlated
run_test = runs.test(fit_stdres, plot.it = TRUE, alternative =
"positive.correlated")
```



run_test

```
##
## Runs Test - Positive Correlated
##
## data: fit_stdres
## Standardized Runs Statistic = -3.062, p-value = 0.001101
```

```
# Compute critical value.
qnorm(0.05, lower.tail = TRUE)
```

```
## [1] -1.645
```

H0: the sequence was produced in a random manner Ha: the sequence was not produced in a random manner

Test statistic: $Z = -3.0615$ Significance level: $\alpha = 0.05$ Critical value (upper tail): $Z_{1-\alpha/2} = 1.96$
Critical region: Reject H0 if $|Z| > 1.96$

Since the test statistic is greater than the critical value ($p\text{-value} < 0.05$) we conclude that the sequence are not random at the 0.05 significance level, indicating error terms in the model are correlated and there is a pattern in the residuals present. This reconfirms earlier conclusion in (a).

Problem 2


```
# Import data
filename = "P229-30.txt"
mydata = read.table(filename, header = T)
```

Part a


```
fit = lm(DJIA ~ Time, mydata)
summary(fit)
```

```
##
## Call:
## lm(formula = DJIA ~ Time, data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -457.3  -111.6    -9.8   145.7   385.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5212.300     22.197   234.8  <2e-16 ***
## Time          4.024       0.146    27.5  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 179 on 260 degrees of freedom
## Multiple R-squared:  0.744,    Adjusted R-squared:  0.743
## F-statistic: 756 on 1 and 260 DF,  p-value: <2e-16
```

```
plot.lm(fit, which = 1) # only get residuals vs fitted
```

```
plot(mydata$Time, mydata$DJIA, xlab = "Time", ylab = "DJIA")  
abline(fit, col = "red")
```



```
library(lmtest)
dwtest(fit, alternative = "two.sided")
```

```
##
## Durbin-watson test
##
## data: fit
## DW = 0.0559, p-value < 2.2e-16
## alternative hypothesis: true autocorrelation is not 0
```

It is not clear what “linear trend model” refers to. If it refers to the linear regression model DJIA vs Time, then the plot DJIA vs Time clearly shows that the linear model is not adequate because of the cyclical behavior. Residual plot shows a trend, suggesting presence of auto-correlation in the residuals and, thus, the linear model does not seem to be adequate as the linear regression assumption of independent-errors does not hold. The presence of correlated errors have an impact on estimates, standard errors and statistical tests.

The graph of residuals show the presence of time dependence in the error term. Autocorrelation might suggest that a time-dependent variable is missing from the model.

The Durbin Watston test (two-sided) suggests that there is autocorrelation present (p-value<0.05)

Part b

```
# Lag functions:
# http://heuristically.wordpress.com/2012/10/29/lag-function-for-data-frames/
# http://ctsinkin.com/2012/03/11/generating-a-laglead-variables/

# Create lag t-1
n = dim(mydata)[1]
mydata_lagged = data.frame(Time_t = mydata$Time[2:n], DJIA_t =
mydata$DJIA[2:n],
  Time_t_1 = mydata$Time[1:n - 1], DJIA_t_1 = mydata$DJIA[1:n - 1])
head(mydata_lagged)
```


```
##      Time_t DJIA_t Time_t_1 DJIA_t_1
## 1         2   5177         1   5117
## 2         3   5194         2   5177
## 3         4   5174         3   5194
## 4         5   5181         4   5174
## 5         6   5198         5   5181
## 6         7   5130         6   5198
```

```
# Regress DJIA t vs DJIA t-1


fit_lagged = lm(DJIA_t ~ DJIA_t_1, mydata_lagged)
summary(fit_lagged)
```

```
##
## Call:
## lm(formula = DJIA_t ~ DJIA_t_1, data = mydata_lagged)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -176.88  -22.40   -0.64    26.48   125.14
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  36.89838   42.98910    0.86    0.39
## DJIA_t_1      0.99446    0.00748  133.00 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 42.4 on 259 degrees of freedom
## Multiple R-squared:  0.986,    Adjusted R-squared:  0.986
## F-statistic: 1.77e+04 on 1 and 259 DF,  p-value: <2e-16
```

```
plot.lm(fit_lagged, which = 1) # only get residuals vs fitted
```



```
plot(mydata_lagged$DJIA_t_1, mydata_lagged$DJIA_t, xlab = "DJIA (t-1)",
     ylab = "DJIA (t)")
abline(fit_lagged, col = "red")
```



```
library(lmtest)
dwtest(fit_lagged, alternative = "two.sided")
```

```
##
## Durbin-Watson test
##
## data: fit_lagged
## DW = 1.759, p-value = 0.04345
## alternative hypothesis: true autocorrelation is not 0
```

The plot DJIA (t) vs DJIA (t-1) shows the linear model might be adequate

The residuals vs Fitted now appears not to show a trend, so there is no strong evidence of autocorrelation in the residuals, indicating that assumption of uncorrelated residuals (independent-errors assumption) might not be violated.


The Durbin Watson test (two-sided is more conservative) suggests that there still might be some autocorrelation (p-value borderline < 0.05), but there is not strong evidence for it. Compared to (a), it is clear that this model is more adequate for a linear regression.

Part c


```
fit = lm(log10(DJIA) ~ Time, mydata)
summary(fit)
```

```
##
## Call:
## lm(formula = log10(DJIA) ~ Time, data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.03481 -0.00842 -0.00012  0.01111  0.02648
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.72e+00   1.64e-03  2271.4  <2e-16 ***
## Time         3.00e-04   1.08e-05    27.8  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0132 on 260 degrees of freedom
## Multiple R-squared:  0.749,    Adjusted R-squared:  0.748
## F-statistic: 775 on 1 and 260 DF,  p-value: <2e-16
```

```
plot.lm(fit, which = 1) # only get residuals vs fitted
```



```
plot(mydata$Time, log10(mydata$DJIA), xlab = "Time", ylab =
"log10(DJIA)")
abline(fit, col = "red")
```



```
library(lmtest)
dwtest(fit, alternative = "two.sided")
```

```
##
## Durbin-Watson test
##
## data: fit
## DW = 0.0601, p-value < 2.2e-16
## alternative hypothesis: true autocorrelation is not 0
```

Note, the problem does not specify what base to use for logarithm, so use base 10.

It is not clear what “linear trend model” refers to. If it refers to the linear regression model DJIA vs Time, then the plot DJIA vs Time clearly shows that the linear model is not adequate because of the cyclical behavior. Residual plot shows a trend, suggesting presence of auto-correlation in the residuals and, thus, the linear model does not seem to be adequate as the linear regression assumption of independent-errors does not hold. The presence of correlated errors have an impact on estimates, standard errors and statistical tests.

The graph of residuals show the presence of time dependence in the error term. Autocorrelation might suggest that a time-dependent variable is missing from the model.

The Durbin Watson test (two-sided) suggests that there is autocorrelation present (p-value<0.05)


```
# Lag functions:
# http://heuristically.wordpress.com/2012/10/29/lag-function-for-data-frames/
# http://ctszkin.com/2012/03/11/generating-a-laglead-variables/

# Regress log10(DJIA t) vs log10(DJIA t-1)


fit_lagged = lm(log10(DJIA_t) ~ log10(DJIA_t_1), mydata_lagged)
summary(fit_lagged)
```

```
##
## Call:
## lm(formula = log10(DJIA_t) ~ log10(DJIA_t_1), data = mydata_lagged)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.013818 -0.001715  0.000003  0.002126  0.008526
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.02698    0.02884    0.94    0.35
## log10(DJIA_t_1) 0.99292    0.00767  129.40 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.00324 on 259 degrees of freedom
## Multiple R-squared:  0.985,    Adjusted R-squared:  0.985
## F-statistic: 1.67e+04 on 1 and 259 DF,  p-value: <2e-16
```

```
plot.lm(fit_lagged, which = 1) # only get residuals vs fitted
```

```
plot(log10(mydata_lagged$DJIA_t_1), log10(mydata_lagged$DJIA_t), xlab =
"log10(DJIA (t-1))",
      ylab = "log10(DJIA (t))")
abline(fit_lagged, col = "red")
```



```
library(lmtest)
dwtest(fit_lagged, alternative = "two.sided")
```

```
##
## Durbin-watson test
##
## data: fit_lagged
## DW = 1.774, p-value = 0.05824
## alternative hypothesis: true autocorrelation is not 0
```

The plot $\log \text{DJIA } (t)$ vs $\log \text{DJIA } (t-1)$ clearly shows the linear model is adequate

The residuals vs Fitted now do not show a trend, so there is no evidence of autocorrelation in the residuals, indicating that assumption of uncorrelated residuals (independent-errors assumption) is not violated.

The Durbin Watston test (two-sided is more conservative) suggests that there is no strong evidence for autocorrelation (p-value borderline > 0.05), although this is borderline. Compared to the previous model, it is clear that this model is more adequate for a linear regression.

The conclusions reached in (a) and (b) are similar. No big differences are noticed. The coefficients estimates change, but the significant tests and R^2 remain almost the same. The plots also show the same patterns. It seems that there is only a change in the scale and decrease in the variability/volatility.

The main difference is the result of the Durbin Watson test, which shows the log model is slightly better than non-log model in reducing autocorrelation, in which the log model now has no strong evidence at a 0.05 significance level for autocorrelation.

The non-log model has the advantage of keeping the same units that is easy to interpret. The log model might be preferred though because of the reduction in variability, symmetrization of the distribution and no strong evidence of autocorrelation.

Problem 3

Part a

```
mydata_lagged$log_DJIA_t = log10(mydata_lagged$DJIA_t)
mydata_lagged$log_DJIA_t_1 = log10(mydata_lagged$DJIA_t_1)
fit_lagged = lm(log_DJIA_t ~ log_DJIA_t_1, mydata_lagged[1:129, ])
summary(fit_lagged)
```

```
##
## Call:
## lm(formula = log_DJIA_t ~ log_DJIA_t_1, data = mydata_lagged[1:129,
##    ])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.013329 -0.001812  0.000058  0.002033  0.008180
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.1587     0.0724    2.19   0.03 *
## log_DJIA_t_1    0.9577     0.0194   49.49 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.00326 on 127 degrees of freedom
## Multiple R-squared:  0.951,    Adjusted R-squared:  0.95
## F-statistic: 2.45e+03 on 1 and 127 DF,  p-value: <2e-16
```

```
# Mean squared error (calcuatue using # obs or df residuals?)
MSE_log = sum((fit_lagged$res)^2)/fit_lagged$df.res
MSE_log
```

```
## [1] 1.064e-05
```

```
# or summary(fit_lagged)$sigma^2 anova(fit_lagged)['Residuals', 'Mean
Sq']
```

For this question, answers are given for non-log model from the previous question for DJIA t vs DJIA t-1 models. See above for discussion on the adequacy of the models. Note the conclusions to the questions below do not change whether log or non-log model is used. Non-log model was chosen because it is simpler and keeps in the same units for the predictions and errors so easier to interpret.

Note that the training set goes from days 1 to 129 for the lagged values. Thus, DJIA(t-1) is for days 1 to 129, and DJIA (t) for days 2 to 130. This ensures only data from the first half of the (130 days) is used. Also note that MSE (using the book definition) is computed using the degrees of freedom of the residuals (not just the mean of the residuals squared).

- The residual mean square (log units): 1.0641×10^{-5}

Part b

First day of July 1996 = day 131, so start with DJIA_t_1 (130)

```
# option 1: use data (not predicted of t becomes t-1 etc)

first_day = 131
last_day = 131 + 15 - 1
newdata = data.frame(log_DJIA_t_1 =
mydata_lagged$log_DJIA_t_1[(first_day -
1):(last_day - 1)])

predicted_log = predict(fit_lagged, newdata)
predicted = 10^predicted_log
actual_log = log10(mydata$DJIA[first_day:last_day])
actual = mydata$DJIA[first_day:last_day]

pred_error_log = actual_log - predicted_log
pred_error = actual - predicted

results = data.frame(date = mydata$Date[first_day:last_day], day =
mydata$Time[first_day:last_day],
  actual = actual, predicted = predicted, pred_error = pred_error,
  actual_log = actual_log,
  predicted_log = predicted_log, pred_error_log = pred_error_log)

results
```

```
##      date day actual predicted pred_error actual_log predicted_log
## 1  7/1/96 131  5730    5653    76.646    3.758    3.752
## 2  7/2/96 132  5720    5725    -5.079    3.757    3.758
## 3  7/3/96 133  5703    5716   -13.252    3.756    3.757
## 4  7/4/96 134  5703    5700    3.362    3.756    3.756
## 5  7/5/96 135  5588    5700  -111.518    3.747    3.756
## 6  7/8/96 136  5551    5590   -38.827    3.744    3.747
## 7  7/9/96 137  5582    5554   27.949    3.747    3.745
## 8  7/10/96 138  5604    5584   20.009    3.748    3.747
## 9  7/11/96 139  5520    5605   -84.014    3.742    3.749
## 10 7/12/96 140  5511    5525   -14.285    3.741    3.742
## 11 7/15/96 141  5350    5515  -165.808    3.728    3.742
## 12 7/16/96 142  5359    5361    -2.094    3.729    3.729
## 13 7/17/96 143  5377    5370    7.149    3.731    3.730
## 14 7/18/96 144  5464    5387   77.062    3.738    3.731
## 15 7/19/96 145  5427    5471  -44.034    3.735    3.738
##      pred_error_log
## 1      0.0058484
## 2     -0.0003854
## 3     -0.0010080
## 4      0.0002561
## 5     -0.0085815
## 6     -0.0030272
## 7      0.0021800
## 8      0.0015535
## 9     -0.0065595
## 10     -0.0011244
## 11     -0.0132565
## 12     -0.0001696
## 13      0.0005779
## 14      0.0061685
## 15     -0.0035097
```

```
# plot(results$day,results$actual) lines(results$day,results$actual,
# col='red') par(new=TRUE)
#
plot(results$day,results$predicted,xlab='',ylab='',ylim=range(results$actual))
# lines(results$day,results$predicted, col='blue')

require(ggplot2)

plot1 = ggplot(results, aes(day)) + geom_point(aes(y = actual), size =
3, color = "red") +
  geom_line(aes(y = actual), colour = "red") + geom_point(aes(y =
predicted),
  size = 3, color = "blue") + geom_line(aes(y = predicted), colour =
"blue") +
  scale_colour_manual("Legend", breaks = c("Actual", "Predicted"),
values = c("red",
  "blue")) + ylab("Actual (red) vs Predicted (blue)")
```

Part c

```
### Part c
AVE_Sq_error15 = mean(pred_error^2)
AVE_Sq_error15
```

```
## [1] 4260
```

```
AVE_Sq_error15_log = mean(pred_error_log^2)
AVE_Sq_error15_log
```

```
## [1] 2.641e-05
```

- Average of the squared error (log units) = 2.6411×10^{-5}
- Average of the squared error (original units) = 4260.0055

As expected, average squared prediction errors are much higher than MSE in (a) since part © is testing data in a new period and in a smaller sample, while MSE (a) is for the data that the model was built on and over a longer time period.

Part d

First day of July 1996 = day 131, so start with DJIA_t_1 (130)

```
# Use to predict second half (132 days) First day of July 1996 = day
131,
# so start with DJIA_t_1 (130)

# option 1: use data

first_day = 131
last_day = dim(mydata)[1]

newdata = data.frame(log_DJIA_t_1 =
mydata_lagged$log_DJIA_t_1[(first_day -
1):(last_day - 1)])

predicted_log = predict(fit_lagged, newdata)
predicted = 10^predicted_log
actual_log = log10(mydata$DJIA[first_day:last_day])
actual = mydata$DJIA[first_day:last_day]

pred_error_log = actual_log - predicted_log
pred_error = actual - predicted

results = data.frame(date = mydata$Date[first_day:last_day], day =
mydata$Time[first_day:last_day],
  actual = actual, predicted = predicted, pred_error = pred_error,
  actual_log = actual_log,
  predicted_log = predicted_log, pred_error_log = pred_error_log)

results
```

##	date	day	actual	predicted	pred_error	actual_log	predicted_log
## 1	7/1/96	131	5730	5653	76.64567	3.758	3.752
## 2	7/2/96	132	5720	5725	-5.07885	3.757	3.758
## 3	7/3/96	133	5703	5716	-13.25204	3.756	3.757

## 4	7/4/96	134	5703	5700	3.36244	3.756	3.756
## 5	7/5/96	135	5588	5700	-111.51756	3.747	3.756
## 6	7/8/96	136	5551	5590	-38.82677	3.744	3.747
## 7	7/9/96	137	5582	5554	27.94909	3.747	3.745
## 8	7/10/96	138	5604	5584	20.00925	3.748	3.747
## 9	7/11/96	139	5520	5605	-84.01357	3.742	3.749
## 10	7/12/96	140	5511	5525	-14.28494	3.741	3.742
## 11	7/15/96	141	5350	5515	-165.80773	3.728	3.742
## 12	7/16/96	142	5359	5361	-2.09351	3.729	3.729
## 13	7/17/96	143	5377	5370	7.14948	3.731	3.730
## 14	7/18/96	144	5464	5387	77.06200	3.738	3.731
## 15	7/19/96	145	5427	5471	-44.03417	3.735	3.738
## 16	7/22/96	146	5391	5435	-44.08629	3.732	3.735
## 17	7/23/96	147	5347	5401	-54.05789	3.728	3.732
## 18	7/24/96	148	5355	5358	-3.32273	3.729	3.729
## 19	7/25/96	149	5422	5366	56.18528	3.734	3.730
## 20	7/26/96	150	5473	5430	42.64721	3.738	3.735
## 21	7/29/96	151	5435	5479	-44.77848	3.735	3.739
## 22	7/30/96	152	5482	5442	39.45150	3.739	3.736
## 23	7/31/96	153	5529	5488	41.03739	3.743	3.739
## 24	8/1/96	154	5595	5533	61.84488	3.748	3.743
## 25	8/2/96	155	5680	5596	83.84140	3.754	3.748
## 26	8/5/96	156	5674	5677	-3.18010	3.754	3.754
## 27	8/6/96	157	5696	5672	23.96292	3.756	3.754
## 28	8/7/96	158	5719	5693	25.62629	3.757	3.755
## 29	8/8/96	159	5713	5715	-1.14557	3.757	3.757
## 30	8/9/96	160	5681	5710	-28.36818	3.754	3.757
## 31	8/12/96	161	5705	5679	26.10313	3.756	3.754
## 32	8/13/96	162	5647	5702	-54.25349	3.752	3.756
## 33	8/14/96	163	5667	5646	20.58321	3.753	3.752
## 34	8/15/96	164	5666	5665	0.71730	3.753	3.753
## 35	8/16/96	165	5689	5664	25.44041	3.755	3.753
## 36	8/19/96	166	5699	5687	12.77117	3.756	3.755
## 37	8/20/96	167	5721	5696	25.02897	3.757	3.756
## 38	8/21/96	168	5690	5717	-27.29419	3.755	3.757
## 39	8/22/96	169	5733	5687	46.44701	3.758	3.755
## 40	8/23/96	170	5723	5729	-6.05848	3.758	3.758
## 41	8/26/96	171	5694	5719	-24.64052	3.755	3.757
## 42	8/27/96	172	5711	5691	20.35122	3.757	3.755
## 43	8/28/96	173	5712	5708	4.82647	3.757	3.756
## 44	8/29/96	174	5648	5709	-60.96586	3.752	3.757
## 45	8/30/96	175	5616	5647	-30.44107	3.749	3.752
## 46	9/2/96	176	5616	5617	-0.33333	3.749	3.749
## 47	9/3/96	177	5648	5617	31.84667	3.752	3.749
## 48	9/4/96	178	5657	5647	9.54037	3.753	3.752
## 49	9/5/96	179	5607	5656	-48.54775	3.749	3.752
## 50	9/6/96	180	5660	5608	52.17606	3.753	3.749
## 51	9/9/96	181	5734	5658	75.49824	3.758	3.753
## 52	9/10/96	182	5727	5729	-1.97253	3.758	3.758
## 53	9/11/96	183	5755	5723	32.14057	3.760	3.758
## 54	9/12/96	184	5772	5749	22.61761	3.761	3.760
## 55	9/13/96	185	5839	5766	72.91474	3.766	3.761
## 56	9/16/96	186	5889	5829	59.91777	3.770	3.766
## 57	9/17/96	187	5889	5878	11.09814	3.770	3.769
## 58	9/18/96	188	5877	5877	-0.01821	3.769	3.769
## 59	9/19/96	189	5868	5866	1.32549	3.768	3.768
## 60	9/20/96	190	5888	5857	31.24154	3.770	3.768
## 61	9/23/96	191	5895	5877	17.71544	3.770	3.769
## 62	9/24/96	192	5874	5883	-8.99698	3.769	3.770
## 63	9/25/96	193	5877	5863	14.12866	3.769	3.768
## 64	9/26/96	194	5869	5866	2.43549	3.769	3.768
## 65	9/27/96	195	5873	5858	14.64043	3.769	3.768
## 66	9/30/96	196	5882	5862	19.99974	3.770	3.768
## 67	10/1/96	197	5905	5871	33.88770	3.771	3.769

## 68	10/2/96	198	5934	5893	41.23266	3.773	3.770
## 69	10/3/96	199	5933	5921	12.33309	3.773	3.772
## 70	10/4/96	200	5993	5919	73.41327	3.778	3.772
## 71	10/7/96	201	5980	5977	3.03486	3.777	3.776
## 72	10/8/96	202	5967	5964	2.45963	3.776	3.776
## 73	10/9/96	203	5931	5952	-21.23399	3.773	3.775
## 74	10/10/96	204	5922	5917	4.35409	3.772	3.772
## 75	10/11/96	205	5969	5909	60.61638	3.776	3.771
## 76	10/14/96	206	6010	5954	55.65273	3.779	3.775
## 77	10/15/96	207	6005	5993	11.63525	3.778	3.778
## 78	10/16/96	208	6021	5988	32.65041	3.780	3.777
## 79	10/17/96	209	6059	6003	55.73213	3.782	3.778
## 80	10/18/96	210	6094	6040	54.10757	3.785	3.781
## 81	10/21/96	211	6091	6074	17.30970	3.785	3.783
## 82	10/22/96	212	6062	6070	-8.55337	3.783	3.783
## 83	10/23/96	213	6036	6043	-6.14454	3.781	3.781
## 84	10/24/96	214	5992	6018	-25.93159	3.778	3.779
## 85	10/25/96	215	6007	5976	30.60780	3.779	3.776
## 86	10/28/96	216	5973	5990	-17.56884	3.776	3.777
## 87	10/29/96	217	6007	5958	49.47262	3.779	3.775
## 88	10/30/96	218	5993	5990	2.93116	3.778	3.777
## 89	10/31/96	219	6029	5977	52.25147	3.780	3.776
## 90	11/1/96	220	6022	6012	10.27869	3.780	3.779
## 91	11/4/96	221	6042	6005	37.14262	3.781	3.778
## 92	11/5/96	222	6081	6023	57.78435	3.784	3.780
## 93	11/6/96	223	6178	6061	116.60562	3.791	3.783
## 94	11/7/96	224	6206	6153	52.82658	3.793	3.789
## 95	11/8/96	225	6220	6180	39.58566	3.794	3.791
## 96	11/11/96	226	6256	6193	62.22430	3.796	3.792
## 97	11/12/96	227	6266	6227	38.54829	3.797	3.794
## 98	11/13/96	228	6274	6237	36.79537	3.798	3.795
## 99	11/14/96	229	6313	6245	67.73843	3.800	3.796
## 100	11/15/96	230	6348	6282	65.82496	3.803	3.798
## 101	11/18/96	231	6347	6316	31.32493	3.803	3.800
## 102	11/19/96	232	6398	6315	83.08206	3.806	3.800
## 103	11/20/96	233	6430	6363	67.21307	3.808	3.804
## 104	11/21/96	234	6418	6394	24.78719	3.807	3.806
## 105	11/22/96	235	6472	6383	89.07633	3.811	3.805
## 106	11/25/96	236	6548	6433	114.36491	3.816	3.808
## 107	11/26/96	237	6528	6506	22.62160	3.815	3.813
## 108	11/27/96	238	6499	6487	11.99357	3.813	3.812
## 109	11/28/96	239	6499	6460	39.66087	3.813	3.810
## 110	11/29/96	240	6522	6460	62.02087	3.814	3.810
## 111	12/2/96	241	6522	6481	40.73933	3.814	3.812
## 112	12/3/96	242	6443	6481	-38.27067	3.809	3.812
## 113	12/4/96	243	6423	6406	17.19244	3.808	3.807
## 114	12/5/96	244	6437	6387	50.15942	3.809	3.805
## 115	12/6/96	245	6382	6400	-18.48472	3.805	3.806
## 116	12/9/96	246	6464	6348	116.04958	3.810	3.803
## 117	12/10/96	247	6473	6426	47.26981	3.811	3.808
## 118	12/11/96	248	6403	6435	-32.32357	3.806	3.809
## 119	12/12/96	249	6304	6367	-63.78302	3.800	3.804
## 120	12/13/96	250	6305	6273	31.51870	3.800	3.797
## 121	12/16/96	251	6268	6274	-6.10686	3.797	3.798
## 122	12/17/96	252	6308	6240	68.68324	3.800	3.795
## 123	12/18/96	253	6347	6278	69.01558	3.803	3.798
## 124	12/19/96	254	6474	6314	159.25545	3.811	3.800
## 125	12/20/96	255	6484	6435	49.18515	3.812	3.809
## 126	12/23/96	256	6489	6445	43.56202	3.812	3.809
## 127	12/24/96	257	6523	6450	72.99417	3.814	3.810
## 128	12/25/96	258	6523	6482	40.79488	3.814	3.812
## 129	12/26/96	259	6547	6482	64.62488	3.816	3.812
## 130	12/27/96	260	6561	6505	56.17781	3.817	3.813
## 131	12/30/96	261	6549	6518	31.09795	3.816	3.814

##	132	12/31/96	262	6448	6507	-59.02182	3.809	3.813
##		pred_error_log						
##	1							
##	2							
##	3							
##	4							
##	5							
##	6							
##	7							
##	8							
##	9							
##	10							
##	11							
##	12							
##	13							
##	14							
##	15							
##	16							
##	17							
##	18							
##	19							
##	20							
##	21							
##	22							
##	23							
##	24							
##	25							
##	26							
##	27							
##	28							
##	29							
##	30							
##	31							
##	32							
##	33							
##	34							
##	35							
##	36							
##	37							
##	38							
##	39							
##	40							
##	41							
##	42							
##	43							
##	44							
##	45							
##	46							
##	47							
##	48							
##	49							
##	50							
##	51							
##	52							
##	53							
##	54							
##	55							
##	56							
##	57							
##	58							
##	59							
##	60							
##	61							
##	62							

## 63	1.045e-03
## 64	1.803e-04
## 65	1.084e-03
## 66	1.479e-03
## 67	2.500e-03
## 68	3.028e-03
## 69	9.037e-04
## 70	5.353e-03
## 71	2.205e-04
## 72	1.791e-04
## 73	-1.552e-03
## 74	3.194e-04
## 75	4.433e-03
## 76	4.040e-03
## 77	8.423e-04
## 78	2.362e-03
## 79	4.013e-03
## 80	3.873e-03
## 81	1.236e-03
## 82	-6.124e-04
## 83	-4.418e-04
## 84	-1.875e-03
## 85	2.219e-03
## 86	-1.276e-03
## 87	3.592e-03
## 88	2.125e-04
## 89	3.780e-03
## 90	7.419e-04
## 91	2.678e-03
## 92	4.146e-03
## 93	8.276e-03
## 94	3.713e-03
## 95	2.773e-03
## 96	4.342e-03
## 97	2.680e-03
## 98	2.554e-03
## 99	4.685e-03
## 100	4.527e-03
## 101	2.149e-03
## 102	5.677e-03
## 103	4.564e-03
## 104	1.680e-03
## 105	6.019e-03
## 106	7.652e-03
## 107	1.507e-03
## 108	8.022e-04
## 109	2.658e-03
## 110	4.150e-03
## 111	2.721e-03
## 112	-2.572e-03
## 113	1.164e-03
## 114	3.397e-03
## 115	-1.256e-03
## 116	7.868e-03
## 117	3.183e-03
## 118	-2.187e-03
## 119	-4.372e-03
## 120	2.177e-03
## 121	-4.229e-04
## 122	4.754e-03
## 123	4.748e-03
## 124	1.082e-02
## 125	3.307e-03
## 126	2.925e-03

```
## 127      4.887e-03
## 128      2.725e-03
## 129      4.308e-03
## 130      3.735e-03
## 131      2.067e-03
## 132     -3.957e-03
```

```
# plot(results$day,results$actual) lines(results$day,results$actual,
# col='red') par(new=TRUE)
#
plot(results$day,results$predicted,xlab='',ylab='',ylim=range(results$actual))
# lines(results$day,results$predicted, col='blue')
require(ggplot2)

plot2 = ggplot(results, aes(day)) + geom_point(aes(y = actual), size =
3, color = "red") +
  geom_line(aes(y = actual), colour = "red") + geom_point(aes(y =
predicted),
  size = 3, color = "blue") + geom_line(aes(y = predicted), colour =
"blue") +
  scale_colour_manual("Legend", breaks = c("Actual", "Predicted"),
values = c("red",
  "blue")) + ylab("Actual (red) vs Predicted (blue)")

AVE_Sq_error132 = mean(pred_error^2)
AVE_Sq_error132
```

```
## [1] 2467
```

```
AVE_Sq_error132_log = mean(pred_error_log^2)
AVE_Sq_error132_log
```

```
## [1] 1.303e-05
```

- Average of the squared error (log units) = 1.3034×10^{-5}
- Average of the squared error (original units) = 2466.9671

Average squared prediction errors are higher than MSE (a), but lower than average squared prediction error for first 15 days of second half of year (part c)

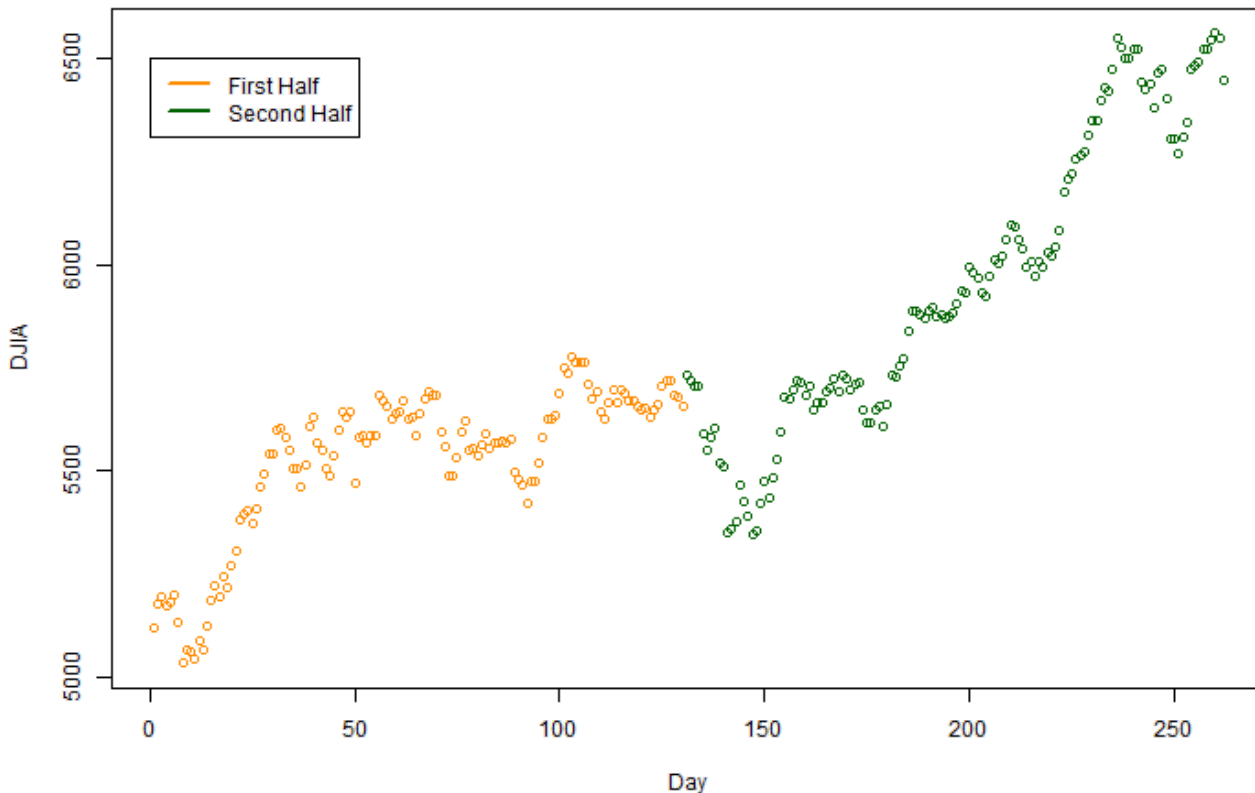
Part e

```

mydata$color[mydata$Time<131]="darkorange"
mydata$color[mydata$Time>=131]="darkgreen"

plot(mydata$Time,mydata$DJIA, xlab="Day", ylab="DJIA", col=mydata$color)
legend(0,6500, c("First Half","Second Half"), # puts text in the legend
in the appropriate place
      lty=c(1,1), # gives the legend appropriate symbols (lines)
      lwd=c(2.5,2.5), # gives the legend lines the correct color and
width
      col=c("darkorange","darkgreen"))

```

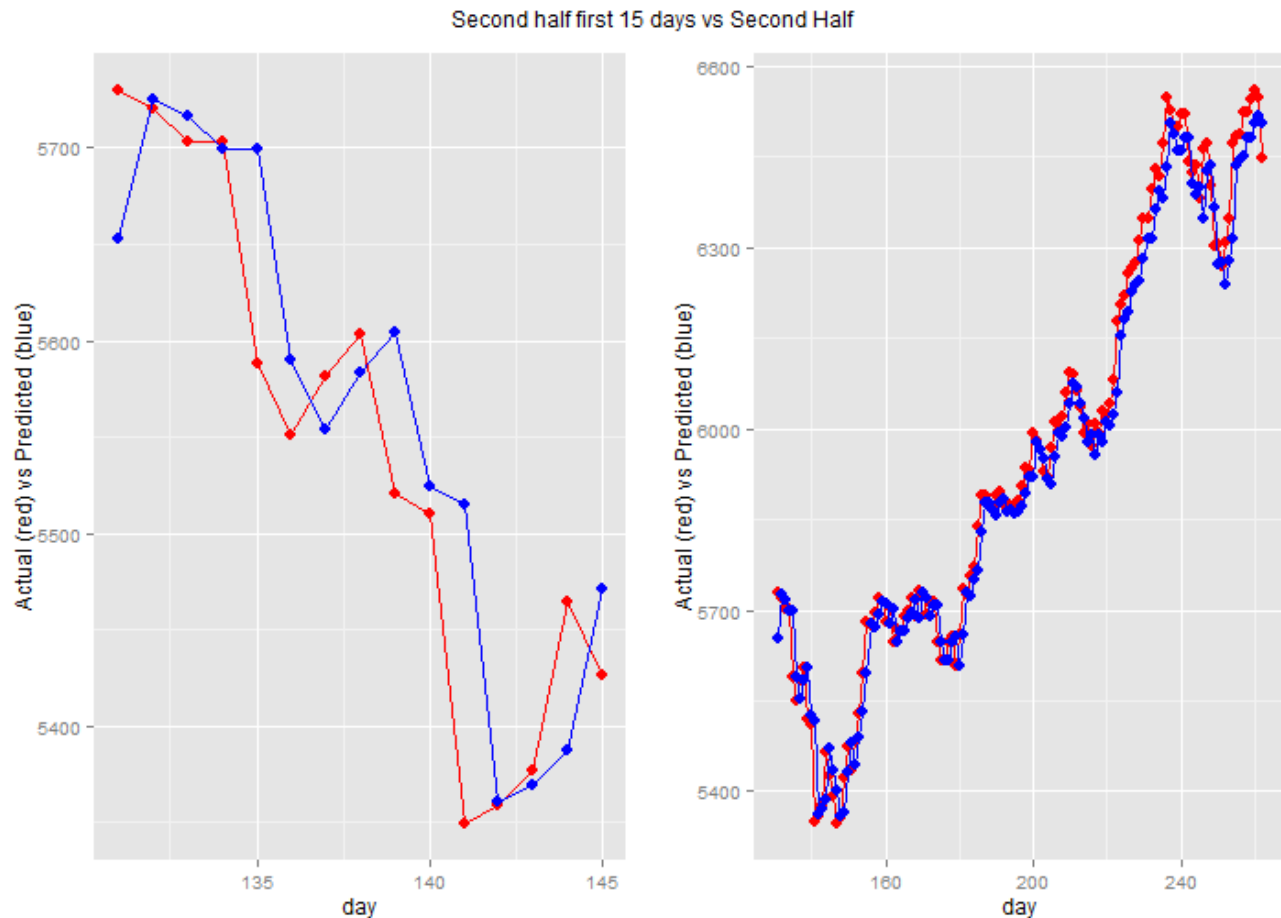


```

require(ggplot2)
require(grid)
require(gridExtra)

grid.arrange(plot1, plot2, ncol=2, main = "second half first 15 days vs
Second Half")

```



From the scatter plot we clearly see a difference between the first half of the year and the second half. Because the model used the training data for the first half of the year, then the prediction error is larger in (c) and (d) than (a) because (c) and (d) are based on the second part of the year, which is different data than the training set. The expected error the model exhibits on new data will always be higher than that it exhibits on the training data (<http://scott.fortmann-roe.com/docs/MeasuringError.html>)

Now usually one would expect that the error in a closer time to the training to be smaller than the error in a further time out. However, this is not the case here because the prediction for period t is a function of the previous period $t-1$. One can explain the results of the average squared prediction error in 15 days © being larger than average squared prediction error for the entire second half of the year (d) as follows:

- The behavior for the entire second half of the year is similar to the behavior of the entire first half of the year. Because the model was based on the entire first half of the year, then the predictions errors will be smaller for the entire second half of the year rather than a small portion (e.g. 15 days). In addition we see that for the first 15 days the DJIA decreases while it increases in a stable manner afterwards, similar to the first half of the year.
- The error of a small sample is also larger than a bigger sample. When we look at the prediction vs actual in the first 15 days, we see day-to-day big changes in the actual values for some days. Thus, since the prediction of the next day is based on the previous period, then the error will be substantial. It is as if the prediction is “catching up” the actual value. In the second half of the year, we see that this happens too, but there are more days in which the day-to-day changes are small, which translates on an smaller prediction average error for the entire second half of the year vs the first 15 days.

Problem 4

```
# Import data
filename = "P329.txt"
mydata = read.table(filename, header = T)

var_names = c("taxes", "bath", "lot", "living", "garage", "rooms",
              "bedrooms",
              "age", "fireplaces", "sale")

predictor_names = c("intercept", "taxes", "bath", "lot", "living",
                    "garage",
                    "rooms", "bedrooms", "age", "fireplaces")

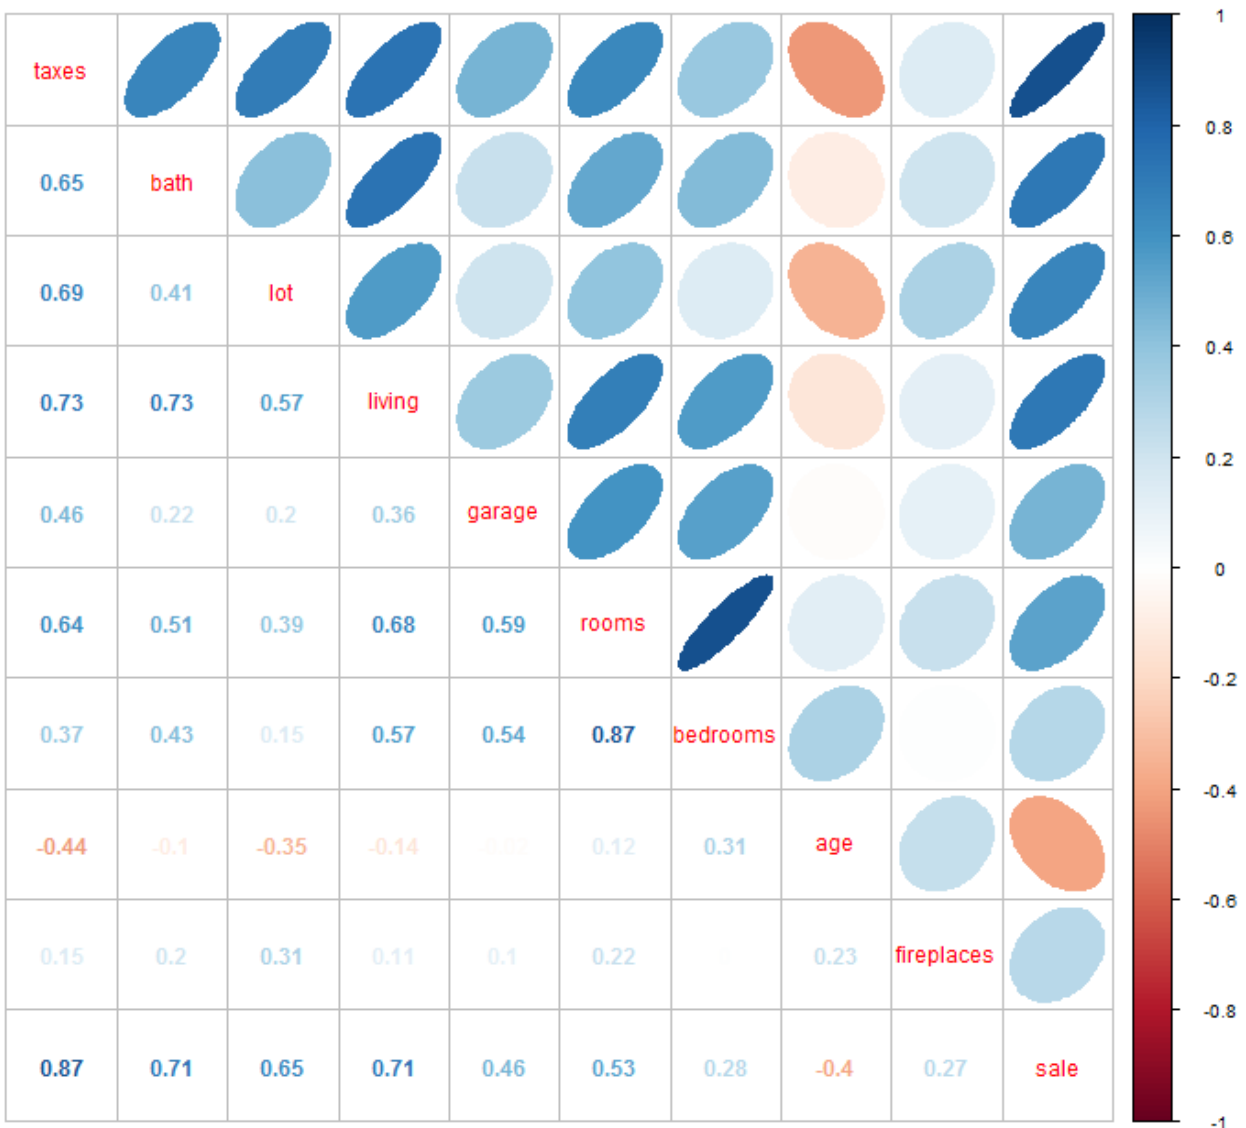
colnames(mydata) = var_names
```

Part a

```
# Correlation matrix

corr = round(cor(mydata), 2)

library(corrplot)
corrplot.mixed(corr, upper = "ellipse", lower = "number")
```



```
# pairs(mydata[,-1], main = 'Correlation coefficients matrix and
scatter
# plot', pch = 21, lower.panel = NULL, panel = panel.smooth,
cex.labels=2)
```

The pairwise correlation coefficients of the predictor variables and the corresponding scatter plots show strong linear relationships among some pairs of predictors variables, suggesting collinearity. (look at high magnitudes for correlation coefficient in conjunction for a trend in the scatter plot)

In particular, rooms (X6) and bedrooms (X7) are strongly correlated, which makes sense since a bedroom is a room too. Thus, both variables cannot be in the model since might cause the non-collinearity assumption to be violated.

```
fit = lm(sale ~ ., mydata)
summary(fit)
```

```
##
## Call:
## lm(formula = sale ~ ., data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.773 -1.980 -0.087  1.662  4.262
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  15.3104     5.9609   2.57    0.022 *
## taxes         1.9541     1.0383   1.88    0.081 .
## bath          6.8455     4.3353   1.58    0.137
## lot           0.1376     0.4944   0.28    0.785
## living        2.7814     4.3948   0.63    0.537
## garage        2.0508     1.3846   1.48    0.161
## rooms        -0.5559     2.3979  -0.23    0.820
## bedrooms     -1.2452     3.4229  -0.36    0.721
## age          -0.0380     0.0673  -0.57    0.581
## fireplaces    1.7045     1.9532   0.87    0.398
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.97 on 14 degrees of freedom
## Multiple R-squared:  0.851,    Adjusted R-squared:  0.756
## F-statistic:  8.9 on 9 and 14 DF,  p-value: 0.000202
```

```
# Compute VIF
library(car)
vif(fit)
```

```
##      taxes      bath      lot      living      garage      rooms
##      7.022      2.835      2.455      3.836      1.823      11.711
## bedrooms      age fireplaces
##      9.722      2.321      1.942
```

```
# Determine VIF > 10
names(vif(fit))[vif(fit) > 10]
```

```
## [1] "rooms"
```

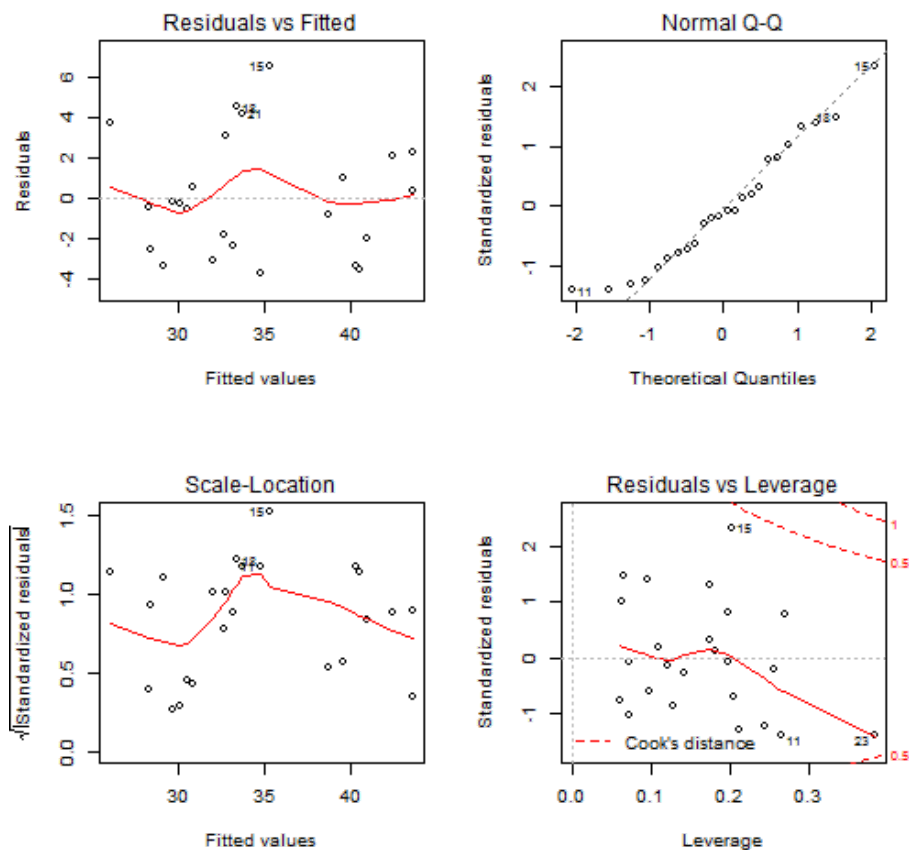
Fitting a linear model with all predictors and computing VIF confirms our suspicion. It appears that rooms (X6) is affected by the presence of collinearity because $VIF > 10$. Thus, there is a multicollinearity problem. Do not include all of them because of multicollinearity. In addition, if all variables in the model are included, none of the variables are significant ($p\text{-value} > 0.05$)

Part b


```
fit = lm(sale ~ taxes + rooms + age, mydata)
summary(fit)
```

```
##
## Call:
## lm(formula = sale ~ taxes + rooms + age, data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.749 -2.408 -0.359  2.138  6.535
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  14.79601    4.97110   2.98  0.00746 **
## taxes         3.48946    0.72937   4.78  0.00011 ***
## rooms        -0.41551    1.18226  -0.35  0.72892
## age           0.00492    0.06360   0.08  0.93906
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.12 on 20 degrees of freedom
## Multiple R-squared:  0.765,    Adjusted R-squared:  0.73
## F-statistic: 21.8 on 3 and 20 DF,  p-value: 1.65e-06
```

```
par(mfrow = c(2, 2))
plot(fit)
```



```
par(mfrow = c(1, 1))

# Compute VIF
library(car)
vif(fit)
```

```
## taxes rooms age
## 3.140 2.580 1.881
```

```
# Determine VIF > 10
names(vif(fit))[vif(fit) > 10]
```

```
## character(0)
```

The residuals diagnostics and VIF show no problems. The R^2 of 0.77 is OK but not great, and the two predictors in the model (rooms and age) are far from being significant (p-value > 0.05). Thus, this model would NOT adequately describe the sale price.

Part c

```
# Stepwise regression to determine best model, start with all variables
library(MASS)
```

```
##
## Attaching package: 'MASS'
##
## The following object(s) are masked from 'package:VGAM':
##
##      huber
```

```
fit = lm(sale ~ ., data = mydata)
fit_stepAIC = step(object = fit, direction = "both") # AIC
```

```
## Start: AIC=59.36
## sale ~ taxes + bath + lot + living + garage + rooms + bedrooms +
## age + fireplaces
##
##           Df Sum of Sq RSS  AIC
## - rooms      1      0.47 124 57.5
## - lot         1      0.68 124 57.5
## - bedrooms    1      1.17 125 57.6
## - age         1      2.82 127 57.9
## - living      1      3.54 127 58.0
## - fireplaces  1      6.73 130 58.6
## <none>                124 59.4
## - garage        1     19.39 143 60.9
## - bath          1     22.04 146 61.3
## - taxes         1     31.30 155 62.8
##
## Step: AIC=57.45
## sale ~ taxes + bath + lot + living + garage + bedrooms + age +
## fireplaces
```

```

##
##          Df Sum of Sq RSS   AIC
## - lot      1      0.8 125 55.6
## - age      1      2.9 127 56.0
## - living   1      3.4 128 56.1
## - fireplaces 1      6.7 131 56.7
## - bedrooms 1      9.4 134 57.2
## <none>      124 57.5
## - garage   1     19.7 144 59.0
## + rooms    1      0.5 124 59.4
## - bath     1     26.2 150 60.0
## - taxes    1     40.2 164 62.2
##
## Step:   AIC=55.61
## sale ~ taxes + bath + living + garage + bedrooms + age + fireplaces
##
##          Df Sum of Sq RSS   AIC
## - age      1      3.4 128 54.2
## - living   1      4.8 130 54.5
## - fireplaces 1      9.6 135 55.4
## - bedrooms 1      9.7 135 55.4
## <none>      125 55.6
## - garage   1     19.0 144 57.0
## + lot      1      0.8 124 57.5
## + rooms    1      0.6 124 57.5
## - bath     1     25.5 150 58.1
## - taxes    1     53.2 178 62.1
##
## Step:   AIC=54.24
## sale ~ taxes + bath + living + garage + bedrooms + fireplaces
##
##          Df Sum of Sq RSS   AIC
## - living   1      5.0 133 53.2
## - fireplaces 1      6.5 135 53.4
## <none>      128 54.2
## + age      1      3.4 125 55.6
## - garage   1     20.0 148 55.7
## + lot      1      1.3 127 56.0
## + rooms    1      0.7 128 56.1
## - bedrooms 1     22.8 151 56.2
## - bath     1     24.3 153 56.4
## - taxes    1     95.1 224 65.6
##
## Step:   AIC=53.16
## sale ~ taxes + bath + garage + bedrooms + fireplaces
##
##          Df Sum of Sq RSS   AIC
## - fireplaces 1      6.2 140 52.3
## <none>      133 53.2
## - garage     1     17.8 151 54.2
## - bedrooms   1     17.9 151 54.2
## + living     1      5.0 128 54.2
## + age        1      3.6 130 54.5
## + lot        1      3.0 130 54.6
## + rooms      1      0.6 133 55.1
## - bath       1     39.3 173 57.4
## - taxes      1    158.0 291 69.9
##
## Step:   AIC=52.26
## sale ~ taxes + bath + garage + bedrooms
##
##          Df Sum of Sq RSS   AIC
## <none>      140 52.3
## + fireplaces 1      6.2 133 53.2

```

```
## + lot          1          5.8 134 53.2
## + living       1          4.7 135 53.4
## - garage       1         20.8 160 53.6
## - bedrooms     1         21.7 161 53.7
## + rooms        1          0.4 139 54.2
## + age          1          0.4 139 54.2
## - bath         1         47.4 187 57.3
## - taxes        1        156.6 296 68.3
```

```
summary(fit_stepAIC)
```

```
##
## Call:
## lm(formula = sale ~ taxes + bath + garage + bedrooms, data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.560 -2.086  0.024  1.858  3.898
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   13.621      3.673    3.71  0.00149 **
## taxes          2.412      0.523    4.62  0.00019 ***
## bath          8.459      3.330    2.54  0.01997 *
## garage        2.060      1.223    1.68  0.10854
## bedrooms     -2.215      1.290   -1.72  0.10218
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.71 on 19 degrees of freedom
## Multiple R-squared:  0.832,    Adjusted R-squared:  0.797
## F-statistic: 23.5 on 4 and 19 DF,  p-value: 3.87e-07
```

```
x = mydata[, 1:9] # design matrix or use model.matrix(fullfit)
y = mydata[, 10] # response vector

# function returns best subset function with different criteria
modelselection = function(x, y) {
  # Inputs: x = design matrix y = response vector
  n = length(y) # number of observations
  p = dim(x)[2] # number of predictors

  # Variable Selection Using Package
  library(leaps)

  # find the best subset
  reg_exh = regsubsets(x, y, nbest = 1, nvmax = n, method =
"exhaustive")
  # summary(reg_exh, matrix.logical=TRUE) names(reg_exh)
  # names(summary(reg_exh))

  # get matrix with models
  models = summary(reg_exh)$which # T/F -> multiply by 1 to get 1/0
(not needed)
  msize = as.numeric(apply(models, 1, sum)) # model size

  # compute criteria
  cp = summary(reg_exh)$cp
  cp = round(cp, 3)
  adjr2 = summary(reg_exh)$adjr2
```

```

adjr2 = round(adjr2, 3)
aic = n * log(summary(reg_exh)$rss/n) + 2 * msize
aic = round(aic, 3)
bic = n * log(summary(reg_exh)$rss/n) + msize * log(n)
bic = round(bic, 3)
# different from regsubsets, just differ by constant bic =
# summary(reg_exh)$bic; bic = round(bic,3)

# alternative optimizing various criteria
leaps(x,y,nbest=1,method='cp')
# leaps(x,y,nbest=1,method='adjr2')

# rank by criteria
rk_cp = as.numeric(factor(cp))
rk_adj2 = vector(length = length(adjr2))
rk_adj2[order(adjr2, decreasing = TRUE)] = 1:length(adjr2) #
highest is better
rk_aic = as.numeric(factor(aic))
rk_bic = as.numeric(factor(bic))

rk_tot = rk_cp + rk_adj2 + rk_aic + rk_bic

# create matrix and data frame of results
results = cbind(msize, models, cp, adjr2, aic, bic, rk_cp, rk_adj2,
rk_aic,
rk_bic, rk_tot)

colnames(results)[2] = "Int"

results_df = data.frame(results)

# display results
results

# alternative x1 = vector(length=length(cp)) x1[order(cp)] =
1:length(cp)

# Models
cp_model = c("intercept", colnames(x)[models[order(cp)[1], ][-1]])
adjr2_model = c("intercept", colnames(x)[models[order(adjr2,
decreasing = TRUE)[1],
][[-1]])
aic_model = c("intercept", colnames(x)[models[order(aic)[1], ][-1]])
bic_model = c("intercept", colnames(x)[models[order(bic)[1], ][-1]])

cat("best cp model:\n", cp_model, "\n")
cat("best adjr2 model:\n", adjr2_model, "\n")
cat("best aic model:\n", aic_model, "\n")
cat("best bic model:\n", bic_model, "\n")

# Order results results[order(cp),]; # order by Cp
# results[order(adjr2,decreasing=TRUE),]; # order by adjr2
# results[order(aic),]; # order by BIC results[order(bic),]; # order
by
# BIC

# alternative sort(cp, decreasing = FALSE,index.return=TRUE)$ix <->
# order(cp)

# plots

plot(reg_exh, scale = "adjr2")
plot(reg_exh, scale = "bic")
plot(reg_exh, scale = "cp")

```

```

localenv = environment()

require(ggplot2)
require(grid)
require(gridExtra)

plot_vector = vector(mode = "list", length = 4)

plot_vector[[1]] = ggplot(results_df, aes(x = results_df[[1]], y =
results_df[[p +
3]]), environment = localenv) + geom_point(size = 4) +
geom_line(aes(y = results_df[[p +
3]]), colour = "blue") + labs(x = colnames(results_df[1]), y =
colnames(results_df[p +
3])) + scale_x_continuous(breaks = msize) + geom_point(data =
results_df[order(cp)[1],
], aes(x = msize, y = cp), colour = "red", size = 5)

plot_vector[[2]] = ggplot(results_df, aes(x = results_df[[1]], y =
results_df[[p +
3 + 1]]), environment = localenv) + geom_point(size = 4) +
geom_line(aes(y = results_df[[p +
3 + 1]]), colour = "blue") + labs(x = colnames(results_df[1]), y
= colnames(results_df[p +
3 + 1])) + scale_x_continuous(breaks = msize) + geom_point(data
= results_df[order(adjr2,
decreasing = TRUE)[1], ], aes(x = msize, y = adjr2), colour =
"red",
size = 5)

plot_vector[[3]] = ggplot(results_df, aes(x = results_df[[1]], y =
results_df[[p +
3 + 2]]), environment = localenv) + geom_point(size = 4) +
geom_line(aes(y = results_df[[p +
3 + 2]]), colour = "blue") + labs(x = colnames(results_df[1]), y
= colnames(results_df[p +
3 + 2])) + scale_x_continuous(breaks = msize) + geom_point(data
= results_df[order(aic)[1],
], aes(x = msize, y = aic), colour = "red", size = 5)

plot_vector[[4]] = ggplot(results_df, aes(x = results_df[[1]], y =
results_df[[p +
3 + 3]]), environment = localenv) + geom_point(size = 4) +
geom_line(aes(y = results_df[[p +
3 + 3]]), colour = "blue") + labs(x = colnames(results_df[1]), y
= colnames(results_df[p +
3 + 3])) + scale_x_continuous(breaks = msize) + geom_point(data
= results_df[order(bic)[1],
], aes(x = msize, y = bic), colour = "red", size = 5)

grid.arrange(plot_vector[[1]], plot_vector[[2]], plot_vector[[3]],
plot_vector[[4]],
ncol = 2, main = "Model Selection")



return(results_df)
}


```

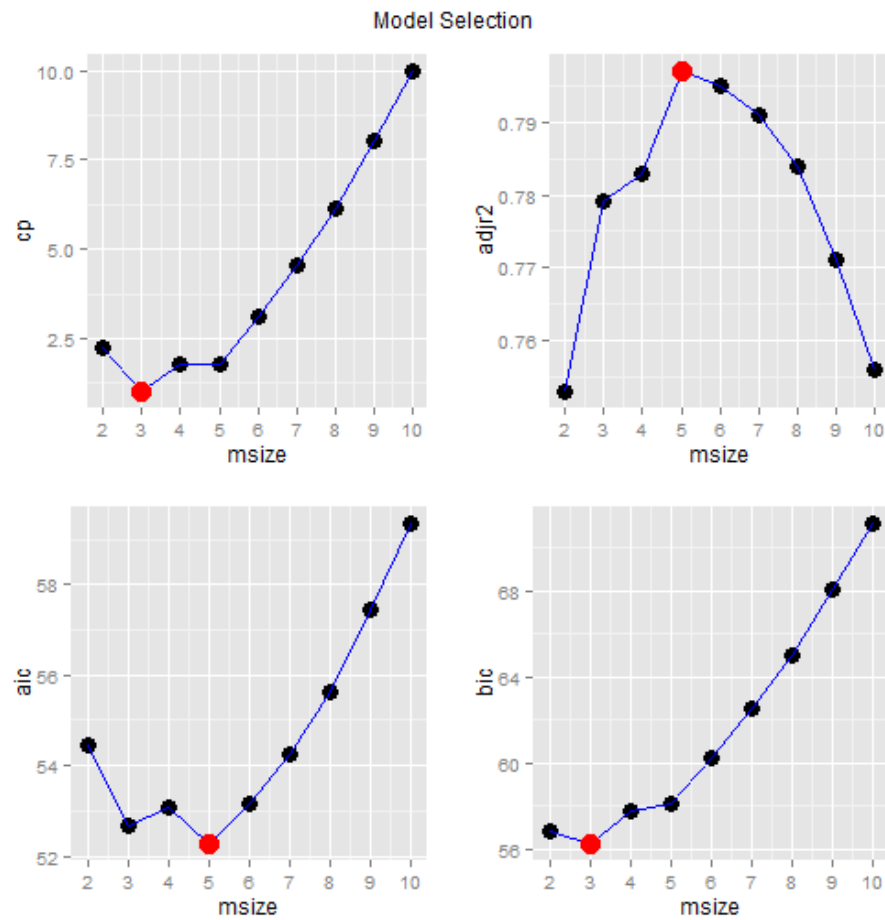
```
bestSubset = modelselection(x, y)
```

```
## warning: package 'leaps' was built under R version 2.15.3
```

```
## best cp model:  
## intercept taxes bath  
## best adjr2 model:  
## intercept taxes bath garage bedrooms  
## best aic model:  
## intercept taxes bath garage bedrooms  
## best bic model:  
## intercept taxes bath
```







bestSubset

```
##      msize Int taxes bath lot living garage rooms bedrooms age
fireplaces
## 1      2    1     1    0  0      0      0      0      0    0
0
## 2      3    1     1    1  0      0      0      0      0    0
0
## 3      4    1     1    1  0      0      0      0      0    0
1
## 4      5    1     1    1  0      0      1      0      1    0
0
## 5      6    1     1    1  0      0      1      0      1    0
1
## 6      7    1     1    1  0      1      1      0      1    0
1
## 7      8    1     1    1  0      1      1      0      1    1
1
## 8      9    1     1    1  1      1      1      0      1    1
1
## 9     10    1     1    1  1      1      1      1      1    1
1
##      cp adjr2    aic    bic rk_cp rk_adjr2 rk_aic rk_bic rk_tot
## 1  2.230 0.753 54.46 56.81      4      9      6      2     21
## 2  0.999 0.779 52.69 56.22      1      6      2      1     10
## 3  1.763 0.783 53.08 57.79      2      5      3      3     13
## 4  1.795 0.797 52.26 58.15      3      1      1      4      9
## 5  3.091 0.795 53.16 60.23      5      2      4      5     16
## 6  4.524 0.791 54.24 62.49      6      3      5      6     20
## 7  6.143 0.784 55.60 65.03      7      4      7      7     25
## 8  8.054 0.771 57.45 68.06      8      7      8      8     31
## 9 10.000 0.756 59.36 71.14      9      8      9      9     35
```

Stepwise regression shows that the best model is: X1 (taxes), X2 (bathrooms), X5 (garage) and X7 (bedrooms).

The model by the expert(sales ~ taxes (x1)) is ranked 4th in cp, last in adjusted r^2 , sixth in aic and 2nd in BIC. Thus, comparing it with other models, the assertion by the expert that the building characteristics are redundant does not hold. For instance, the most adequate models (include intercept) seem to be:

- 1) X1 (taxes) and X2 (bathrooms), -> best in cp and bic criteria
- 2) X1 (taxes), X2 (bathrooms), X5 (garage) and X7 (bedrooms) -> best in adjr^2 and aic criteria

In addition to taxes, the two models include bathroom as a significant predictor of price, and one model includes garage and bedroom as well. Model X1 (taxes), X2 (bathrooms), X5 (garage) and X7 (bedrooms) indicates that garage and bedrooms are not significant. Thus, the most adequate model for predicting sales price seems to be X1 (taxes) and X2 (bathrooms).

Problem 5

```
# Import data
filename = "P256.txt"
mydata = read.table(filename, header = T)

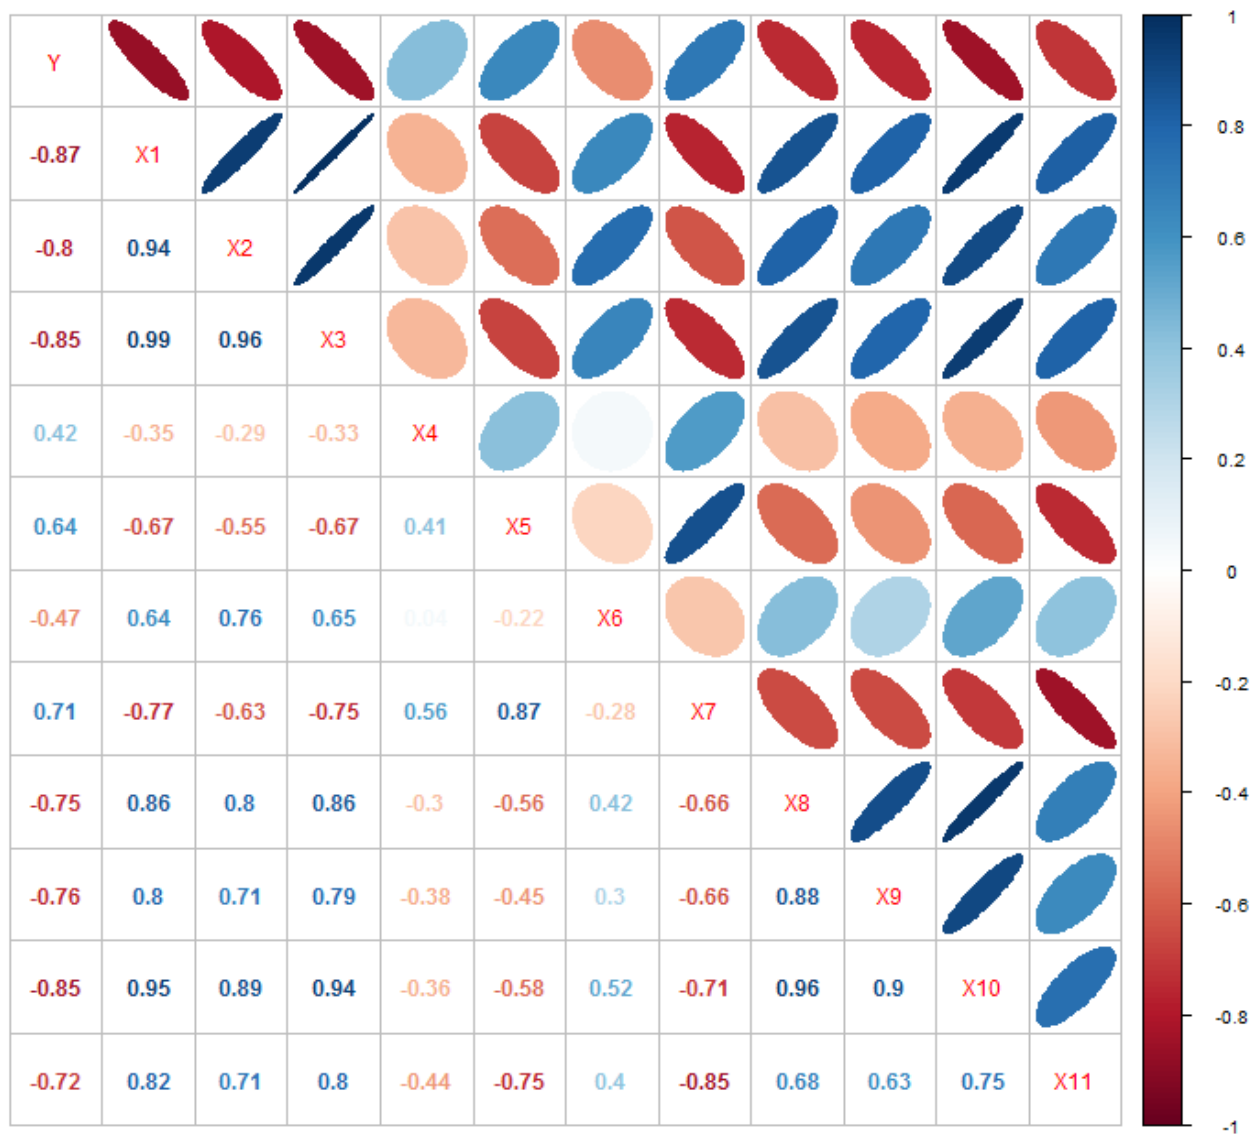
# Look at data
names(mydata)
head(mydata)
nrow(mydata)
summary(mydata)

# Fix names
names(mydata)[-1] = sapply(1:11, function(i) paste("X", i, sep = ""))
```

Part a

```
corr = round(cor(mydata), 2)

library(corrplot)
corrplot.mixed(corr, upper = "ellipse", lower = "number")
```



```
# pairs(mydata[,-1], main = 'Correlation coefficients matrix and
scatter
# plot', pch = 21, lower.panel = NULL, panel = panel.smooth,
cex.labels=2)
```

The pairwise correlation coefficients of the predictor variables and the corresponding scatter plots show strong linear relationships among some pairs of predictor variables, suggesting collinearity. (look at high magnitudes for correlation coefficient in conjunction for a trend in the scatter plot)

For example, X1 is strongly correlated with X2, X3, X8, X9, X10 and X11. If all of these are included in the model, the non-collinearity assumption of the predictors might be violated.

```
fit = lm(Y ~ ., mydata)
summary(fit)
```

```
##
## Call:
## lm(formula = Y ~ ., data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.35  -1.62  -0.60   1.52   5.28
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  17.77320    30.50877   0.58   0.567
## x1          -0.07795     0.05861  -1.33   0.200
## x2          -0.07340     0.08892  -0.83   0.420
## x3           0.12111     0.09135   1.33   0.201
## x4           1.32903     3.09954   0.43   0.673
## x5           5.97599     3.15865   1.89   0.075
## x6           0.30418     1.28909   0.24   0.816
## x7          -3.19858     3.10544  -1.03   0.317
## x8           0.18536     0.12925   1.43   0.169
## x9          -0.39915     0.32381  -1.23   0.234
## x10          -0.00519     0.00589  -0.88   0.390
## x11           0.59865     3.02068   0.20   0.845
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.23 on 18 degrees of freedom
## Multiple R-squared:  0.835,    Adjusted R-squared:  0.735
## F-statistic:  8.3 on 11 and 18 DF,  p-value: 5.29e-05
```

```
# Compute VIF
library(car)
vif(fit)
```

```
##      x1      x2      x3      x4      x5      x6      x7      x8
## 128.835  43.921 160.436   2.058   7.781   5.327  11.735  20.586
## 9.419
##      x10      x11
##  85.676   5.143
```

```
# Determine VIF > 10
names(vif(fit))[vif(fit) > 10]
```

```
## [1] "x1" "x2" "x3" "x7" "x8" "x10"
```

Fitting a linear model with all predictors and computing VIF confirms our suspicion. It appears that X1, X2, X3, X7, X8 and X10 are affected by the presence of collinearity because $VIF > 10$. Thus, there is a multicollinearity problem if all variables are included. Thus, do not include all of them because of multicollinearity. In addition, if all variables in the model are included, none of the variables are significant ($p\text{-value} > 0.05$).

Part b

```

x = mydata[, 2:12] # design matrix or use model.matrix(fullfit)
y = mydata[, 1]   # response vector
# x$X12=x$X2*x$X10 x$X13=x$X8/x$X10

n = length(y)     # number of observations
p = dim(x)[2]     # number of predictors

models = vector(mode = "list", length = 6)
models[[1]] = lm(Y ~ X1, mydata)
models[[2]] = lm(Y ~ X10, mydata)
models[[3]] = lm(Y ~ X1 + X10, mydata)
models[[4]] = lm(Y ~ X2 + X10, mydata)
models[[5]] = lm(Y ~ X8 + X10, mydata)
models[[6]] = lm(Y ~ X8 + X5 + X10, mydata)

full = lm(Y ~ ., mydata)

# compute the selection model criteria input = lm object for desired
# and full model
computeCriteria = function(fit, full) {
  n = length(summary(fit)$res)
  msize = dim(summary(fit)$coeff)[1]
  RSS = sum(summary(fit)$residuals^2)
  cp = RSS/summary(full)$sigma^2 + 2 * msize - n
  adjr2 = summary(fit)$adj.r
  aic = n * log(RSS/n) + 2 * msize
  bic = n * log(RSS/n) + msize * log(n)

  return(round(c(cp, adjr2, aic, bic), 3))
}

results = matrix(, nrow = 6, ncol = 4)

for (i in 1:6) {
  results[i, ] = computeCriteria(models[[i]], full)
}

colnames(results) = c("cp", "adjr2", "aic", "bic")
rownames(results) = 1:6
results

```

```

##      cp adjr2  aic  bic
## 1  0.218 0.751 70.24 73.04
## 2  3.844 0.717 74.12 76.93
## 3  1.660 0.748 71.59 75.80
## 4  5.036 0.715 75.30 79.50
## 5  0.992 0.754 70.80 75.00
## 6 -0.524 0.781 68.25 73.86

```

```

# get best model from six
rownames(results)[order(results[, "cp"])[1]]

```

```
## [1] "6"
```

```
rownames(results)[order(results[, "adjr2"], decreasing = TRUE)[1]]
```

```
## [1] "6"
```

```
rownames(results)[order(results[, "aic"])[1]]
```

```
## [1] "6"
```



```
rownames(results)[order(results[, "bic"])[1]]
```



```
## [1] "1"
```

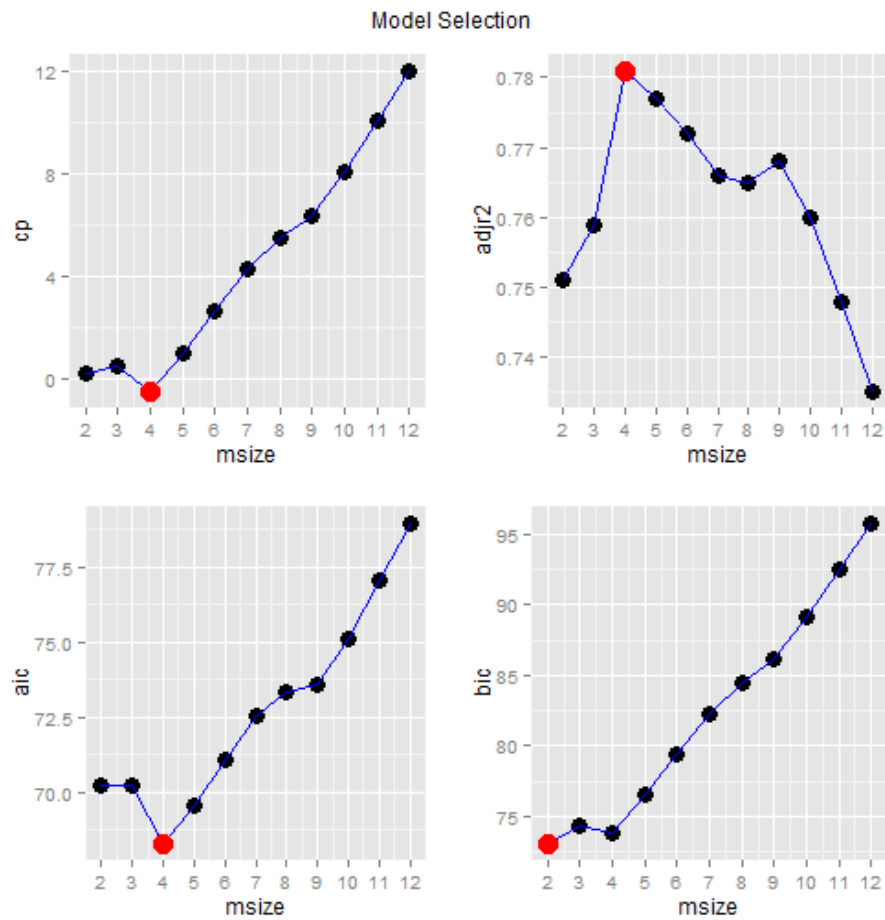
Among the six regression models, model 6 (X8,X5,X10) is the best in predicting Y because it has the highest adjr2, lowest cp, lowest aic and second-to-lowest bic (very close to the lowest bic).

```
# find a better model
bestSubset = modelSelection(x, y)
```

```
## best cp model:
## intercept x5 x8 x10
## best adjr2 model:
## intercept x5 x8 x10
## best aic model:
## intercept x5 x8 x10
## best bic model:
## intercept x1
```





bestSubset

```
##      msize Int  x1 x2 x3 x4 x5 x6 x7 x8 x9 x10 x11      cp adjr2   aic
bic
## 1      2    1  1  0  0  0  0  0  0  0  0  0  0  0.218 0.751 70.24
73.04
## 2      3    1  1  0  0  1  0  0  0  0  0  0  0  0.495 0.759 70.20
74.40
## 3      4    1  0  0  0  0  1  0  0  1  0  1  0 -0.524 0.781 68.25
73.86
## 4      5    1  0  0  0  0  1  0  0  1  1  1  0  0.992 0.777 69.57
76.58
## 5      6    1  0  0  0  0  1  0  1  1  1  1  0  2.633 0.772 71.05
79.46
## 6      7    1  0  0  0  1  1  0  1  1  1  1  0  4.268 0.766 72.52
82.33
## 7      8    1  1  1  1  0  1  0  0  1  1  1  0  5.469 0.765 73.31
84.52
## 8      9    1  1  1  1  0  1  0  1  1  1  1  0  6.362 0.768 73.55
86.17
## 9     10    1  1  1  1  1  1  0  1  1  1  1  0  8.087 0.760 75.10
89.11
## 10    11    1  1  1  1  1  1  1  1  1  1  1  0 10.039 0.748 77.02
92.44
## 11    12    1  1  1  1  1  1  1  1  1  1  1  1 12.000 0.735 78.96
95.77
##      rk_cp rk_adjr2 rk_aic rk_bic rk_tot
## 1      2      9      4      1      16
## 2      3      8      3      3      17
## 3      1      1      1      2      5
## 4      4      2      2      4      12
## 5      5      3      5      5      18
## 6      6      5      6      6      23
## 7      7      6      7      7      27
## 8      8      4      8      8      28
## 9      9      7      9      9      34
## 10     10     10     10     10     40
## 11     11     11     11     11     44
```

Comparing the best models of each model size, we see that the same model (X8,X5,X10) is the best model in terms of adjr2, cp and aic. The bic is also very close to the best model. Thus, no other better models can be suggested (this is if assuming no transformation or higher order terms or interactions terms are considered). If for example an interaction is allowed (e.g consider $x_{12}=x_2 \times x_{10}$), then the best model in terms of cp would be intercept X2 X8 X10 X11 X12. There are many more interactions and transformations (e.g. x_8/x_{10}) that can be tested if one wants to really find the best model.

Stepwise regression to determine best model, start with all variables

```
library(MASS)
fit_stepAIC = step(object = full, direction = "both") # AIC
```

```
## Start: AIC=78.96
## Y ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 + x8 + x9 + x10 + x11
##
##      Df Sum of Sq RSS AIC
## - x11  1      0.4 188 77.0
## - x6   1      0.6 188 77.0
## - x4   1      1.9 189 77.3
## - x2   1      7.1 194 78.1
## - x10  1      8.1 196 78.2
## - x7   1     11.0 198 78.7
## <none>      187 79.0
```

```

## - x9      1      15.8 203 79.4
## - x3      1      18.3 206 79.8
## - x1      1      18.4 206 79.8
## - x8      1      21.4 209 80.2
## - x5      1      37.3 225 82.4
##
## Step: AIC=77.02
## Y ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 + x8 + x9 + x10
##
##      Df Sum of Sq RSS  AIC
## - x6      1      0.5 188 75.1
## - x4      1      1.9 190 75.3
## - x2      1      7.2 195 76.2
## - x10     1      8.2 196 76.3
## <none>      188 77.0
## - x7      1     14.0 202 77.2
## - x9      1     16.9 205 77.6
## - x1      1     18.0 206 77.8
## - x3      1     18.3 206 77.8
## - x8      1     21.6 209 78.3
## + x11     1      0.4 187 79.0
## - x5      1     37.7 226 80.5
##
## Step: AIC=75.1
## Y ~ x1 + x2 + x3 + x4 + x5 + x7 + x8 + x9 + x10
##
##      Df Sum of Sq RSS  AIC
## - x4      1      2.9 191 73.6
## - x2      1      7.9 196 74.3
## - x10     1      9.6 198 74.6
## <none>      188 75.1
## - x7      1     14.4 203 75.3
## - x9      1     18.1 206 75.9
## - x3      1     18.5 207 75.9
## - x1      1     18.6 207 75.9
## - x8      1     22.5 211 76.5
## + x6      1      0.5 188 77.0
## + x11     1      0.3 188 77.0
## - x5      1     40.4 229 78.9
##
## Step: AIC=73.55
## Y ~ x1 + x2 + x3 + x5 + x7 + x8 + x9 + x10
##
##      Df Sum of Sq RSS  AIC
## - x7      1     11.5 203 73.3
## - x2      1     11.6 203 73.3
## <none>      191 73.6
## - x10     1     14.4 206 73.7
## - x1      1     17.0 208 74.1
## - x9      1     18.2 209 74.3
## - x3      1     22.1 213 74.8
## + x4      1      2.9 188 75.1
## + x6      1      1.5 190 75.3
## + x11     1      0.2 191 75.5
## - x8      1     29.3 220 75.8
## - x5      1     41.5 233 77.4
##
## Step: AIC=73.31
## Y ~ x1 + x2 + x3 + x5 + x8 + x9 + x10
##
##      Df Sum of Sq RSS  AIC
## - x1      1      8.7 211 72.6
## - x9      1      9.1 212 72.6
## - x2      1     10.6 213 72.8

```

```

## <none>                203 73.3
## - x3      1          15.1 218 73.5
## + x7      1          11.5 191 73.6
## - x10     1          19.5 222 74.1
## + x11     1           2.9 200 74.9
## - x8      1          27.6 230 75.1
## + x6      1           0.9 202 75.2
## + x4      1           0.0 203 75.3
## - x5      1          34.5 237 76.0
##
## Step: AIC=72.57
## Y ~ x2 + x3 + x5 + x8 + x9 + x10
##
##      Df Sum of Sq  RSS   AIC
## - x9      1      6.1 217 71.4
## - x3      1      6.5 218 71.5
## - x2      1      6.9 218 71.5
## <none>                211 72.6
## + x1      1      8.7 203 73.3
## + x7      1      3.2 208 74.1
## + x11     1      0.5 211 74.5
## + x6      1      0.1 211 74.5
## + x4      1      0.1 211 74.5
## - x5      1     38.9 250 75.6
## - x10     1     54.5 266 77.5
## - x8      1     55.8 267 77.6
##
## Step: AIC=71.42
## Y ~ x2 + x3 + x5 + x8 + x10
##
##      Df Sum of Sq  RSS   AIC
## - x2      1      4.5 222 70.0
## - x3      1      6.1 224 70.3
## <none>                217 71.4
## + x9      1      6.1 211 72.6
## + x1      1      5.7 212 72.6
## + x4      1      0.8 217 73.3
## + x7      1      0.5 217 73.3
## + x11     1      0.5 217 73.3
## + x6      1      0.1 217 73.4
## - x5      1     35.4 253 73.9
## - x8      1     53.7 271 76.0
## - x10     1     87.5 305 79.6
##
## Step: AIC=70.03
## Y ~ x3 + x5 + x8 + x10
##
##      Df Sum of Sq  RSS   AIC
## - x3      1      1.7 224 68.3
## <none>                222 70.0
## + x2      1      4.5 217 71.4
## + x9      1      3.7 218 71.5
## + x1      1      3.6 218 71.5
## + x4      1      1.7 220 71.8
## + x11     1      1.1 221 71.9
## + x7      1      1.1 221 71.9
## + x6      1      0.5 221 72.0
## - x5      1     34.2 256 72.3
## - x8      1     50.4 272 74.2
## - x10     1     84.7 306 77.7
##
## Step: AIC=68.25
## Y ~ x5 + x8 + x10
##

```

```
##          Df Sum of Sq RSS   AIC
## <none>          224 68.3
## + x9           1     5.0 219 69.6
## + x4           1     2.7 221 69.9
## + x3           1     1.7 222 70.0
## + x11          1     1.3 222 70.1
## + x7           1     0.6 223 70.2
## + x6           1     0.1 223 70.2
## + x2           1     0.0 224 70.3
## + x1           1     0.0 224 70.3
## - x5           1    36.6 260 70.8
## - x8           1    53.1 277 72.6
## - x10          1   194.7 418 85.0
```


```
summary(fit_stepAIC)
```

```
##
## Call:
## lm(formula = Y ~ x5 + x8 + x10, data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.593 -1.967 -0.644  2.031  5.882
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.49497    11.76475   0.38    0.706
## x5           2.60734     1.26379   2.06    0.049 *
## x8           0.21812     0.08776   2.49    0.020 *
## x10          -0.00948     0.00199  -4.76   6.4e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.93 on 26 degrees of freedom
## Multiple R-squared:  0.803,    Adjusted R-squared:  0.781
## F-statistic: 35.4 on 3 and 26 DF,  p-value: 2.47e-09
```

Using a stepwise regression confirms the conclusions reached above

Part c

```
plot1 = ggplot(mydata, aes(x = x1, y = Y)) + geom_point(size = 3)
plot2 = ggplot(mydata, aes(x = x2, y = Y)) + geom_point(size = 3)
plot3 = ggplot(mydata, aes(x = x8, y = Y)) + geom_point(size = 3)
plot4 = ggplot(mydata, aes(x = x10, y = Y)) + geom_point(size = 3)
grid.arrange(plot1, plot2, plot3, plot4, ncol = 2, main = "Y vs
x1,x2,x8,x10")
```




The plots suggests that the relationship between Y and X1,X2,X8 and X10 (individually) is not linear. It seems that the relationship is hyperbolic (i.e. $1/x$)

Part d

```
mydata$w = 100/mydata$Y

plot1 = ggplot(mydata, aes(x = x1, y = w)) + geom_point(size = 3)
plot2 = ggplot(mydata, aes(x = x2, y = w)) + geom_point(size = 3)
plot3 = ggplot(mydata, aes(x = x8, y = w)) + geom_point(size = 3)
plot4 = ggplot(mydata, aes(x = x10, y = w)) + geom_point(size = 3)
grid.arrange(plot1, plot2, plot3, plot4, ncol = 2, main = "w vs
x1,x2,x8,x10")
```

The plots now suggests that the relationship between W and X1,X2,X8 and X10 (individually) is more linear than that between Y and the variables.

Part e

```

w = mydata$W # response vector

models = vector(mode = "list", length = 6)
models[[1]] = lm(W ~ X1, mydata)
models[[2]] = lm(W ~ X10, mydata)
models[[3]] = lm(W ~ X1 + X10, mydata)
models[[4]] = lm(W ~ X2 + X10, mydata)
models[[5]] = lm(W ~ X8 + X10, mydata)
models[[6]] = lm(W ~ X8 + X5 + X10, mydata)

full = lm(W ~ ., mydata)

# compute the selection model criteria input = lm object for desired
model
# and full model

results = matrix(, nrow = 6, ncol = 4)

for (i in 1:6) {
  results[i, ] = computeCriteria(models[[i]], full)
}

colnames(results) = c("cp", "adjr2", "aic", "bic")
rownames(results) = 1:6
results

```

```

##      cp adjr2    aic    bic
## 1 156.8 0.743 -15.63 -12.833
## 2 183.9 0.705 -11.49  -8.686
## 3 155.9 0.738 -14.13  -9.923
## 4 154.8 0.740 -14.31 -10.107
## 5 113.2 0.800 -22.25 -18.042
## 6 114.8 0.793 -20.33 -14.725

```

```

# get best model from six
rownames(results)[order(results[, "cp"])[1]]

```

```
## [1] "5"
```

```
rownames(results)[order(results[, "adjr2"], decreasing = TRUE)[1]]
```

```
## [1] "5"
```

```
rownames(results)[order(results[, "aic"])[1]]
```

```
## [1] "5"
```



```
rownames(results)[order(results[, "bic"])[1]]
```


```
## [1] "5"
```

Among the six regression models, model 5 (x8,10) is the best in predicting W because it has the highest adjr2, lowest cp, lowest aic and lowest bic. This answer is different from (b).

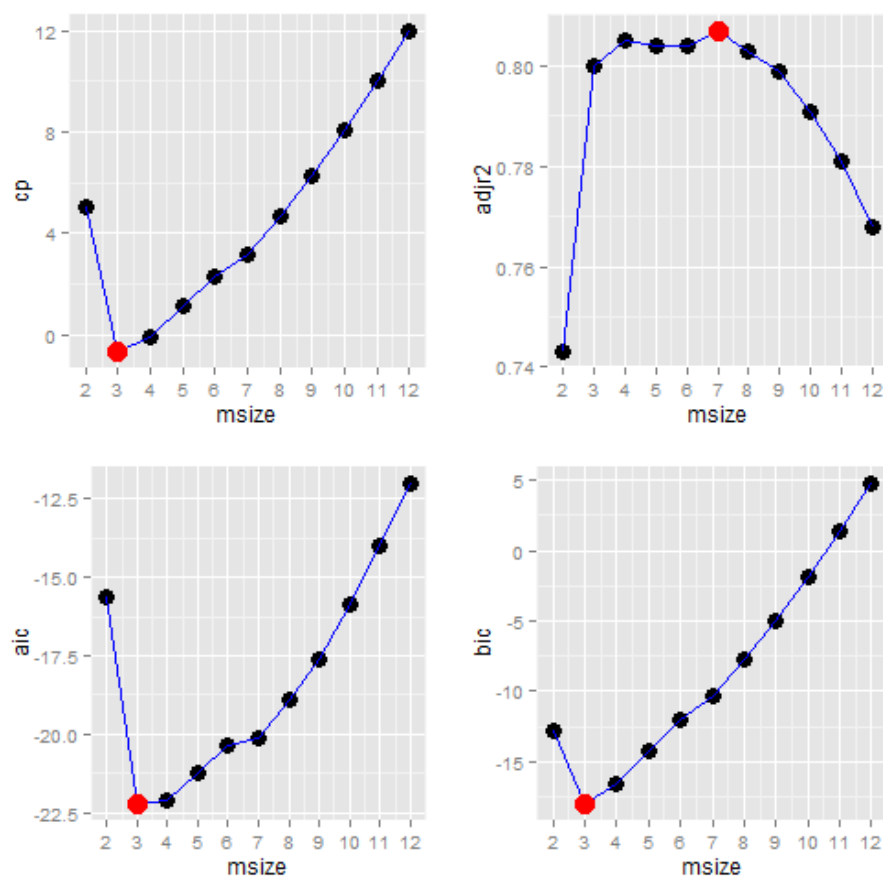
```
# find a better model  
bestSubset = modelSelection(x, w)
```

```
## best cp model:  
## intercept x8 x10  
## best adjr2 model:  
## intercept x1 x2 x3 x6 x8 x10  
## best aic model:  
## intercept x8 x10  
## best bic model:  
## intercept x8 x10
```





Model Selection



bestSubset

```
##      msize Int  x1  x2  x3  x4  x5  x6  x7  x8  x9  x10 x11      cp adjr2      aic
## 1       2    1    1    0    0    0    0    0    0    0    0    0    0    5.051 0.743 -15.63
## 2       3    1    0    0    0    0    0    0    0    1    0    1    0   -0.696 0.800 -22.25
## 3       4    1    0    1    0    0    0    0    0    1    0    1    0   -0.112 0.805 -22.13
## 4       5    1    0    1    1    0    0    0    0    1    0    1    0    1.109 0.804 -21.21
## 5       6    1    0    1    1    0    0    0    1    1    0    1    0    2.311 0.804 -20.37
## 6       7    1    1    1    1    0    0    1    0    1    0    1    0    3.159 0.807 -20.12
## 7       8    1    1    1    1    0    0    1    0    1    0    1    1    4.682 0.803 -18.88
## 8       9    1    1    1    1    0    0    1    1    1    0    1    1    6.231 0.799 -17.61
## 9      10    1    1    1    1    0    1    1    1    1    0    1    1    8.076 0.791 -15.87
## 10     11    1    1    1    1    0    1    1    1    1    1    1    1   10.002 0.781 -13.99
## 11     12    1    1    1    1    1    1    1    1    1    1    1    1   12.000 0.768 -11.99
##      bic rk_cp rk_adjr2 rk_aic rk_bic rk_tot
## 1   -12.833      7      11      9      4     31
## 2   -18.042      1       6      1      1      9
## 3   -16.521      2       2      2      2      8
## 4   -14.207      3       3      3      3     12
## 5   -11.963      4       4      4      5     17
## 6   -10.313      5       1      5      6     17
## 7    -7.667      6       5      6      7     24
## 8    -5.000      8       7      7      8     30
## 9    -1.854      9       8      8      9     34
## 10    1.423     10      9     10     10     39
## 11    4.821     11     10     11     11     43
```

Comparing the best models of each model size, we see that the same model (X8,X10) is the best model in terms of bic, cp and aic. The adjr2 is also very close to the best. Thus, no other better models can be suggested (assuming no transformation or higher order terms or interactions terms are considered).

```
# Stepwise regression to determine best model, start with all variables
library(MASS)
fit_stepAIC = step(object = full, direction = "both") # AIC
```

```
## Start: AIC=-64.9
## W ~ Y + X1 + X2 + X3 + X4 + X5 + X6 + X7 + X8 + X9 + X10 + X11
##
##      Df Sum of Sq  RSS   AIC
## - x3    1      0.00  1.45 -66.9
## - x1    1      0.02  1.47 -66.6
## - x4    1      0.06  1.51 -65.7
## <none>          1.45 -64.9
## - x8    1      0.16  1.61 -63.7
## - x10   1      0.24  1.69 -62.4
## - x11   1      0.29  1.74 -61.5
## - x6    1      0.33  1.78 -60.7
## - x2    1      0.33  1.78 -60.7
## - x9    1      0.91  2.36 -52.3
## - x7    1      1.34  2.79 -47.3
## - x5    1      2.04  3.49 -40.6
## - Y     1      7.59  9.04 -12.0
##
## Step: AIC=-66.88
## W ~ Y + X1 + X2 + X4 + X5 + X6 + X7 + X8 + X9 + X10 + X11
##
##      Df Sum of Sq  RSS   AIC
## - x1    1      0.03  1.48 -68.3
## - x4    1      0.07  1.52 -67.5
## <none>          1.45 -66.9
## - x8    1      0.17  1.62 -65.6
## + x3    1      0.00  1.45 -64.9
## - x10   1      0.24  1.69 -64.4
## - x11   1      0.29  1.74 -63.5
## - x6    1      0.39  1.84 -61.8
## - x2    1      0.94  2.39 -53.9
## - x9    1      0.97  2.42 -53.6
## - x7    1      1.39  2.84 -48.8
## - x5    1      2.57  4.02 -38.3
## - Y     1      8.29  9.74 -11.8
##
## Step: AIC=-68.33
## W ~ Y + X2 + X4 + X5 + X6 + X7 + X8 + X9 + X10 + X11
##
##      Df Sum of Sq  RSS   AIC
## - x4    1      0.05  1.53 -69.4
## <none>          1.48 -68.3
## - x8    1      0.14  1.62 -67.6
## + x1    1      0.03  1.45 -66.9
## + x3    1      0.01  1.47 -66.6
## - x10   1      0.28  1.75 -65.2
## - x11   1      0.36  1.83 -63.8
## - x6    1      0.46  1.94 -62.2
## - x9    1      0.94  2.42 -55.6
## - x2    1      1.01  2.49 -54.7
## - x7    1      1.40  2.88 -50.3
## - x5    1      2.72  4.20 -39.0
## - Y     1      8.29  9.77 -13.7
```

```
##
## Step: AIC=-69.36
## W ~ Y + X2 + X5 + X6 + X7 + X8 + X9 + X10 + X11
##
##      Df Sum of Sq  RSS   AIC
## <none>            1.53 -69.4
## - X8      1      0.13  1.66 -68.9
## + X4      1      0.05  1.48 -68.3
## + X1      1      0.01  1.52 -67.5
## + X3      1      0.00  1.53 -67.4
## - X10     1      0.28  1.81 -66.3
## - X11     1      0.34  1.87 -65.3
## - X6      1      0.41  1.94 -64.2
## - X9      1      0.91  2.43 -57.3
## - X2      1      0.96  2.49 -56.7
## - X7      1      1.42  2.95 -51.6
## - X5      1      2.71  4.24 -40.7
## - Y       1      8.30  9.82 -15.5
```

```
summary(fit_stepAIC)
```

```
##
## Call:
## lm(formula = W ~ Y + X2 + X5 + X6 + X7 + X8 + X9 + X10 + X11,
##     data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.4722 -0.1843  0.0068  0.1592  0.4520
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  13.942140   1.863459   7.48  3.2e-07 ***
## Y            -0.196935   0.018882  -10.43  1.6e-09 ***
## X2             0.014732   0.004147   3.55  0.00200 **
## X5             1.490785   0.250047   5.96  7.9e-06 ***
## X6            -0.219383   0.094632  -2.32  0.03114 *
## X7            -1.012612   0.234490  -4.32  0.00033 ***
## X8            -0.013644   0.010361  -1.32  0.20277
## X9            -0.094223   0.027306  -3.45  0.00253 **
## X10           0.000693   0.000361   1.92  0.06885 .
## X11          -0.530714   0.249935  -2.12  0.04639 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.276 on 20 degrees of freedom
## Multiple R-squared:  0.976,    Adjusted R-squared:  0.965
## F-statistic: 89.4 on 9 and 20 DF,  p-value: 3.58e-14
```

Using a stepwise regression shows a different model, however this model has insignificant terms and also ranks lower in other criteria and also has too many predictors. So the (X8,X10) model is preferred. In conclusion, the transformation of the variable makes a difference in variable selection so it should be examined carefully.


Part f


```
mydata$X13 = mydata$X8/mydata$X10
```

```
fit = lm(Y ~ X13, mydata)
summary(fit)
```

```
##
## Call:
## lm(formula = Y ~ X13, data = mydata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.713  -1.246  -0.023   1.421   4.346
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -11.4         2.1   -5.41  8.9e-06 ***
## X13             566.0        37.2   15.21  4.6e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.09 on 28 degrees of freedom
## Multiple R-squared:  0.892,    Adjusted R-squared:  0.888
## F-statistic: 231 on 1 and 28 DF,  p-value: 4.59e-15
```

```
ggplot(mydata, aes(x = X13, y = Y)) + geom_point(size = 3) +
  stat_smooth(method = "lm",
    se = FALSE)
```



The model seems to be very good in terms of R^2 and fit to the data.